

ER Model, Relational Model

Mapping ER to Relational

COMP 1531

Aarthi Natarajan

Week 09

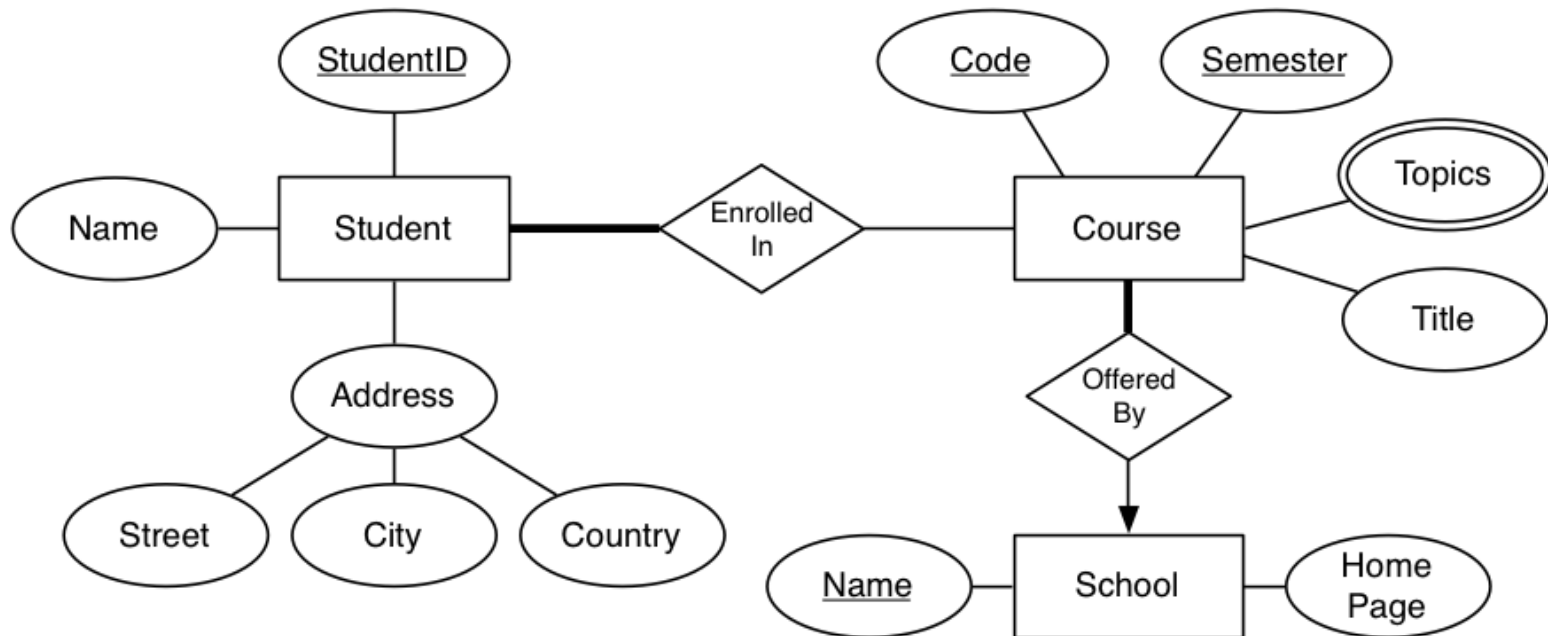
Entity Relationship Diagrams

Story so far...

ER Diagrams

- Last week, we looked at entity-sets and entity instances, attributes, keys and defining relationships (cardinality, degree, participation)
- Putting it all together...

Example 1:

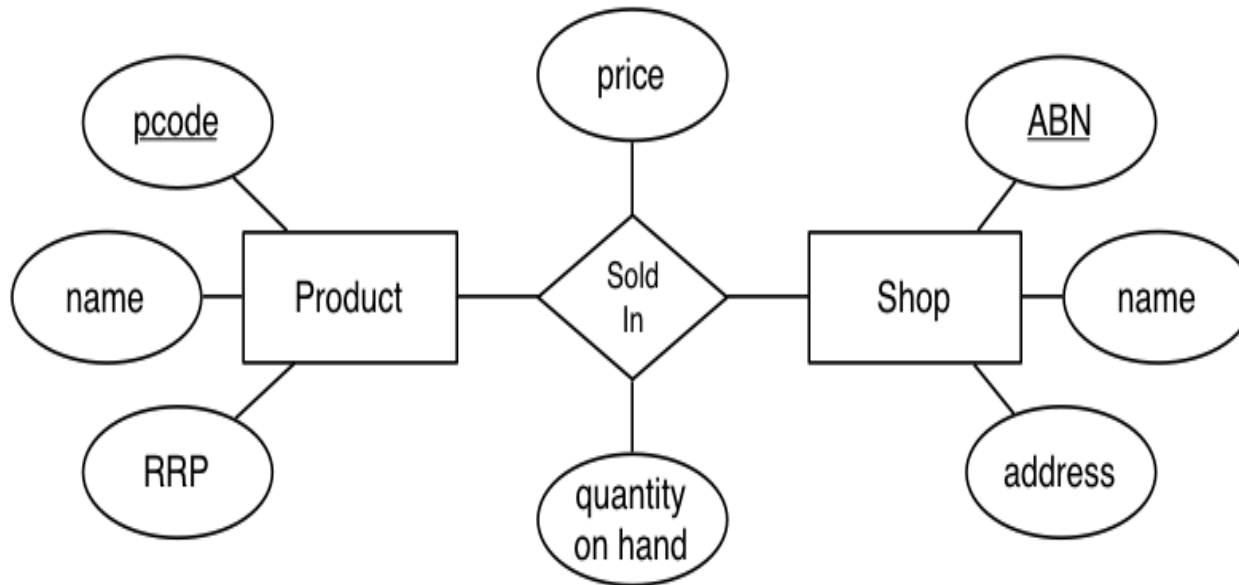


primary key attributes are underlined e.g. StudentID#

Relationship Type with attributes

In some cases, a relationship needs associated attributes

Example:



(price and quantity are related to products in a particular shop)

Summarising notations in an Entity Relationship Diagram


Specific visual symbols indicate different ER design elements:



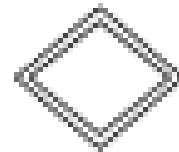
Entity



Weak entity




Relationship



Identifying Relationship



Attribute



Multi-valued attribute

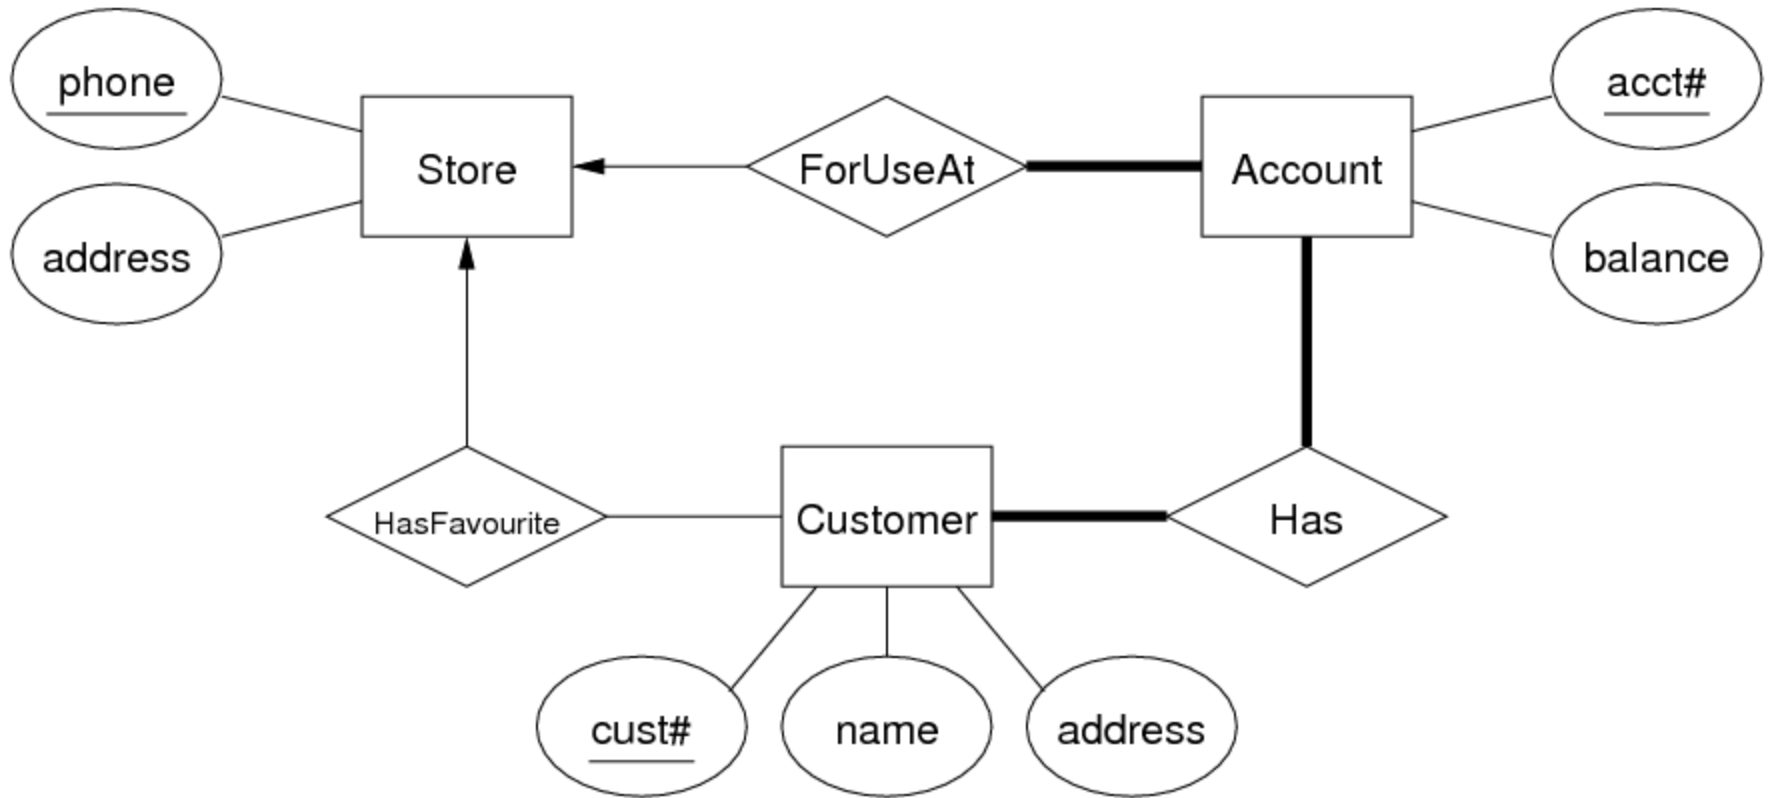


Inheritance



Derived attribute

Another Example:



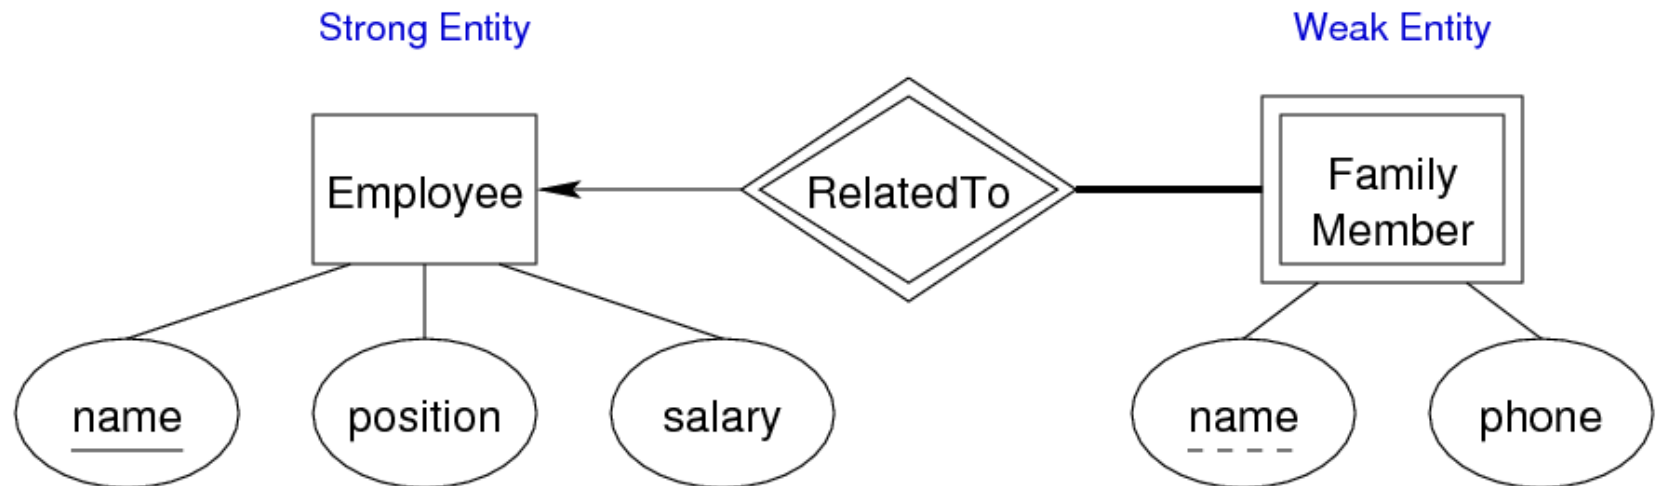
key attributes are underlined e.g. cust#

Weak Entity Set

A Weak entity set

- has no key of its own;
- exist only because of association with strong entities

Example:



ER model vs OO model

Analogy between ER and OO models:

- an **entity** is like an **object instance**
- an **entity set** is like a **class**

Differences between ER and OO models:

- ER modelling doesn't consider operations (methods)

Subclasses and Inheritance

A **subclass** of an entity set A is a set of entities:

- with all attributes of A , plus (usually) it own attributes
- that is involved in all of A 's relationships, plus its own

Properties of subclasses:

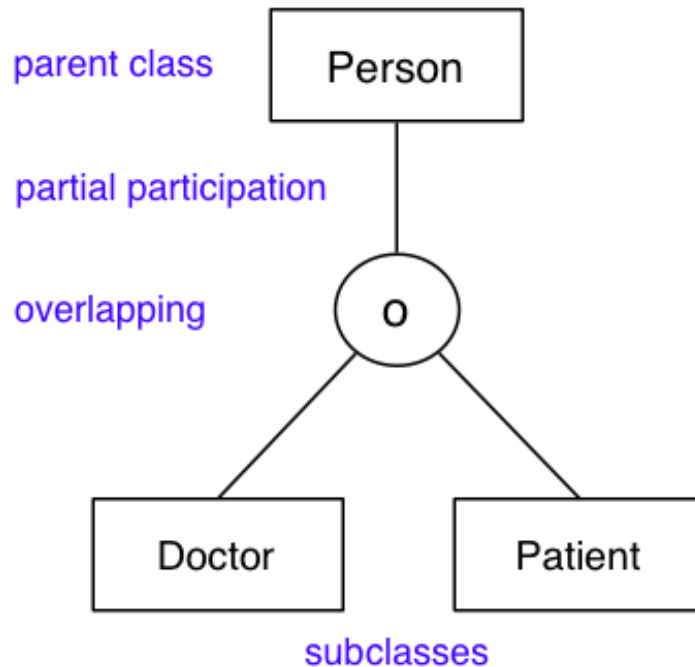
- **overlapping** or **disjoint** (can an entity be in multiple subclasses?)
- **total** or **partial** (does every entity have to also be in a subclass?)

Special case: entity has one subclass ("B is-a A" specialisation)

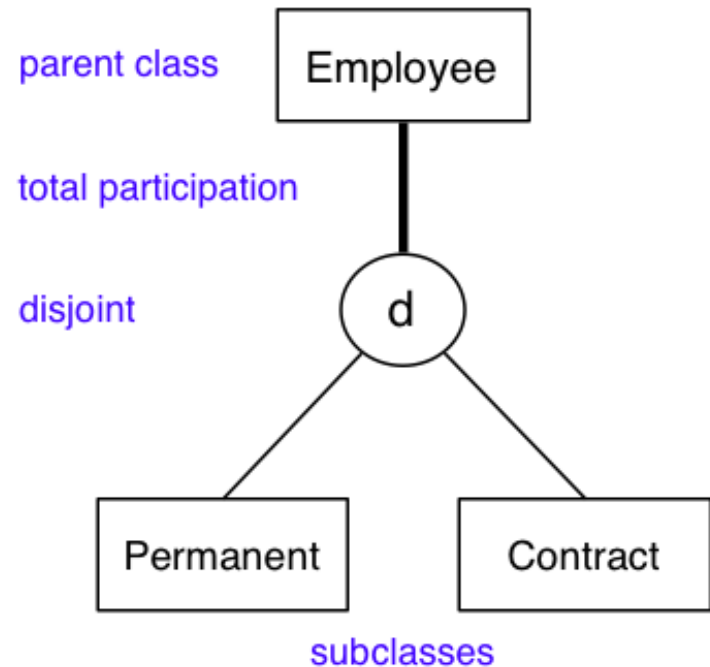
Subclasses and Inheritance

Example:

A person may be a doctor and/or may be a patient or may be neither



Every employee is either a permanent employee or works under a contract



Design considerations using the ER model

- should an "object" be represented by an attribute or entity?
- is a "concept" best expressed as an entity or relationship?
- should we use n -way relⁿship or several 2-way relⁿships?
- is an "object" a strong or weak entity? (usually strong)
- are there subclasses/superclasses within the entities?

Answers to above are worked out by *thinking* about the application domain.

Exercise 1:

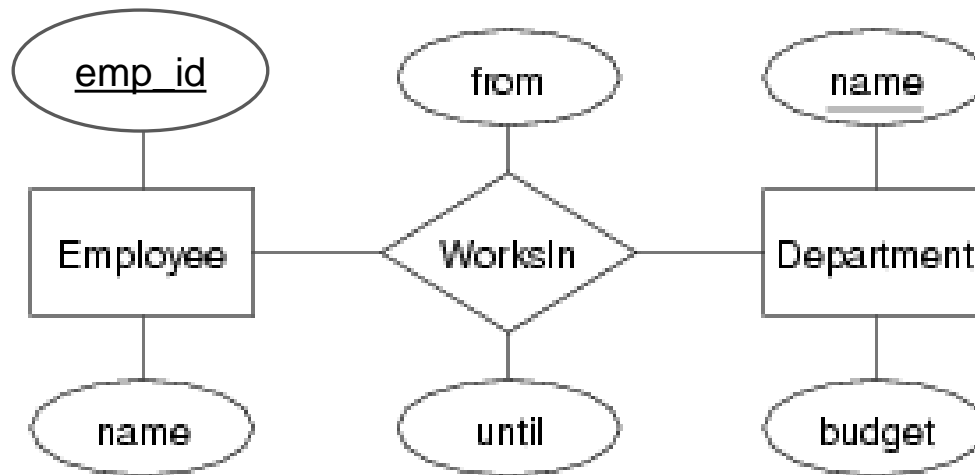
Develop an ER design for the following scenario:

A database records information about employees and the departments they work for:

- For each employee, the name and emp_id
- For each department, the name and allocated budget
- An employee may work for several departments for different periods of time
- A department may have several employees working for it

Design considerations ...

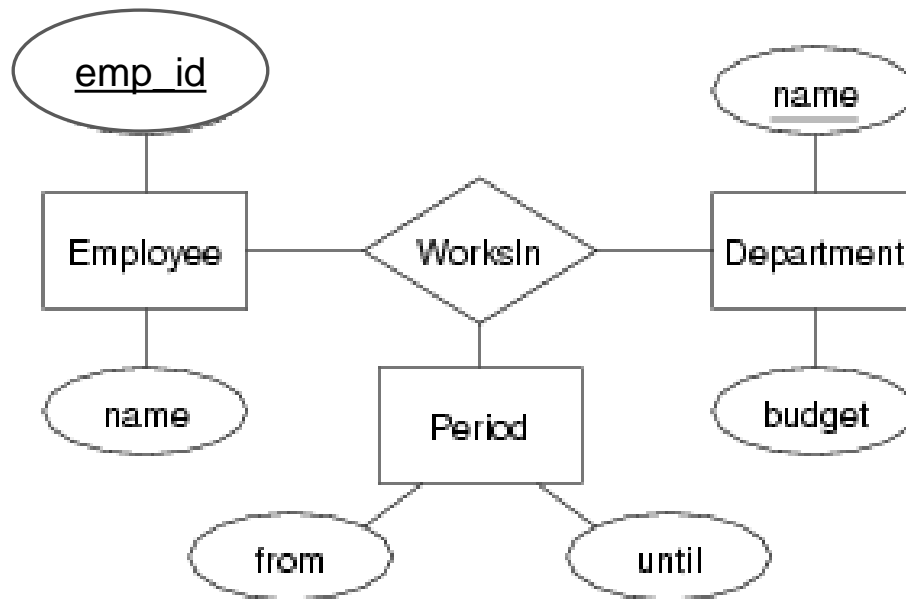
Attribute vs Entity Example (v1)



Assumption: Employees can work for several departments, but cannot work for the same department over two different time periods.

Design considerations ...

Attribute vs Entity Example (v2)



Assumption: Employees can work for the same department over two different time periods.

Design using the ER model

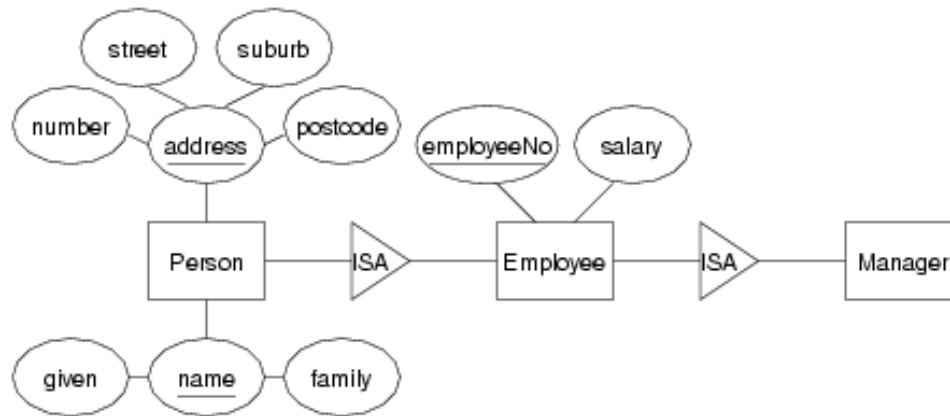
ER diagrams are typically too large to fit on a single screen.
(or a single sheet of paper, if printing)

One commonly used strategy:

- define entity sets separately, showing attributes
- combine entities and relationships on a single diagram (but without showing entity attributes)
- if very large design, may use several linked diagrams as seen in the example in the next three set of slides

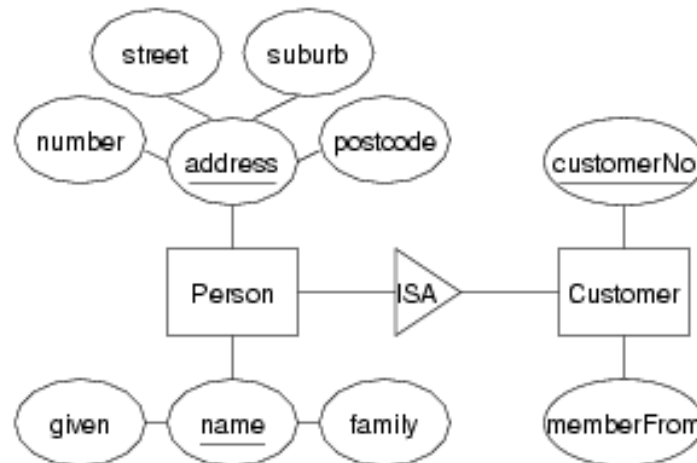
e.g. an ER model for a Bank

(1) Modelling people (employees)



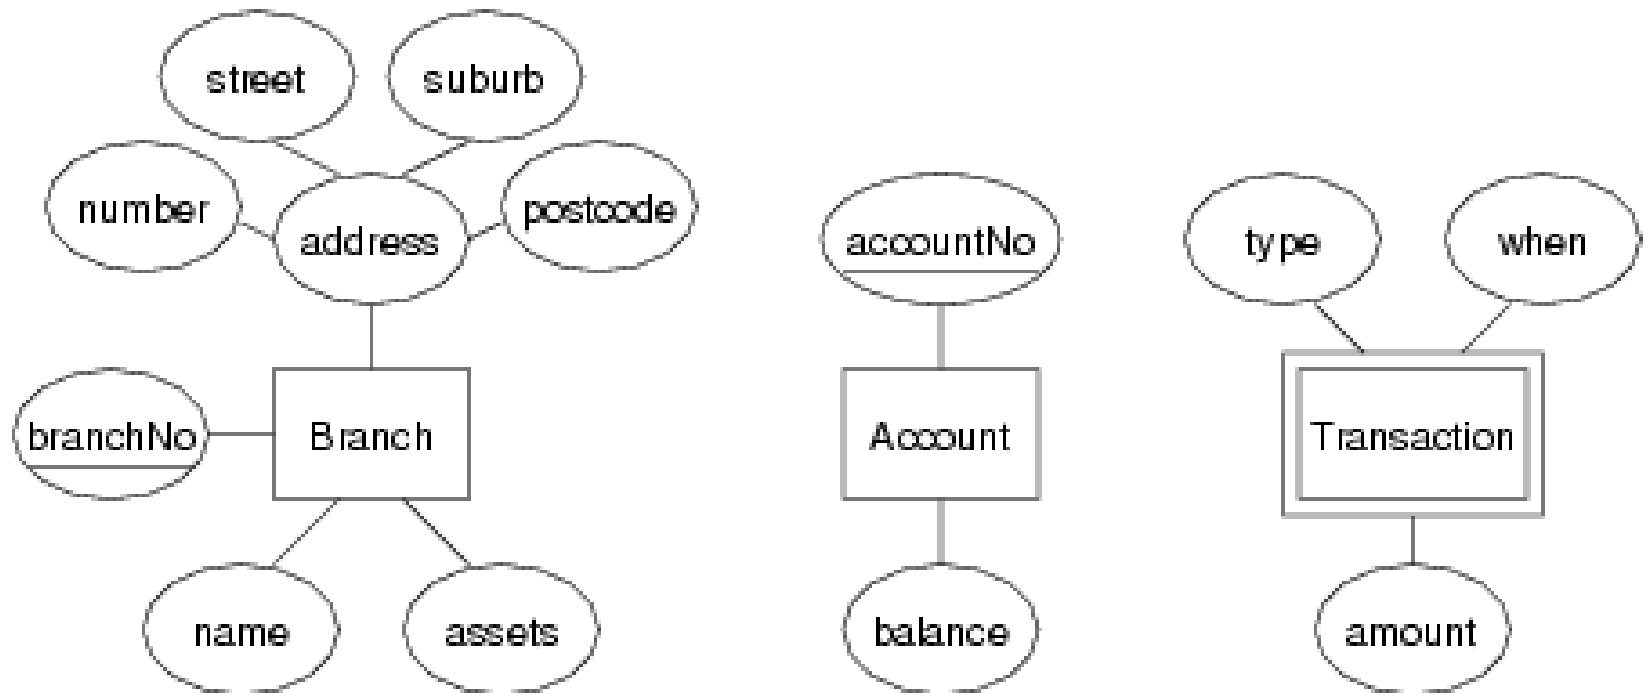
(2) Modelling people (customer)

Modelling people (cont):



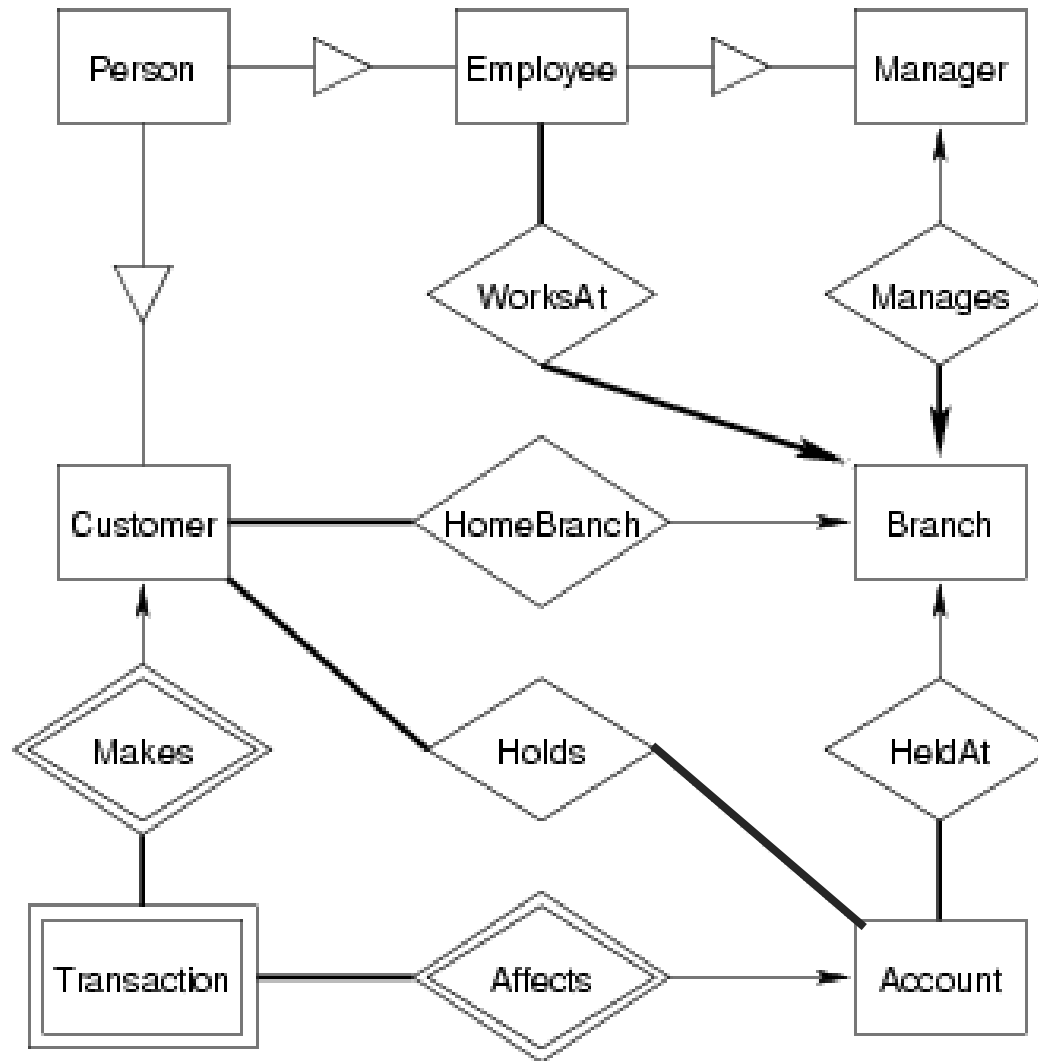
e.g. an ER model for a Bank

(3) Modelling branches, accounts, transactions



e.g. an ER model for a Bank

(4) Putting it all together with relationships




Weak Entity Set

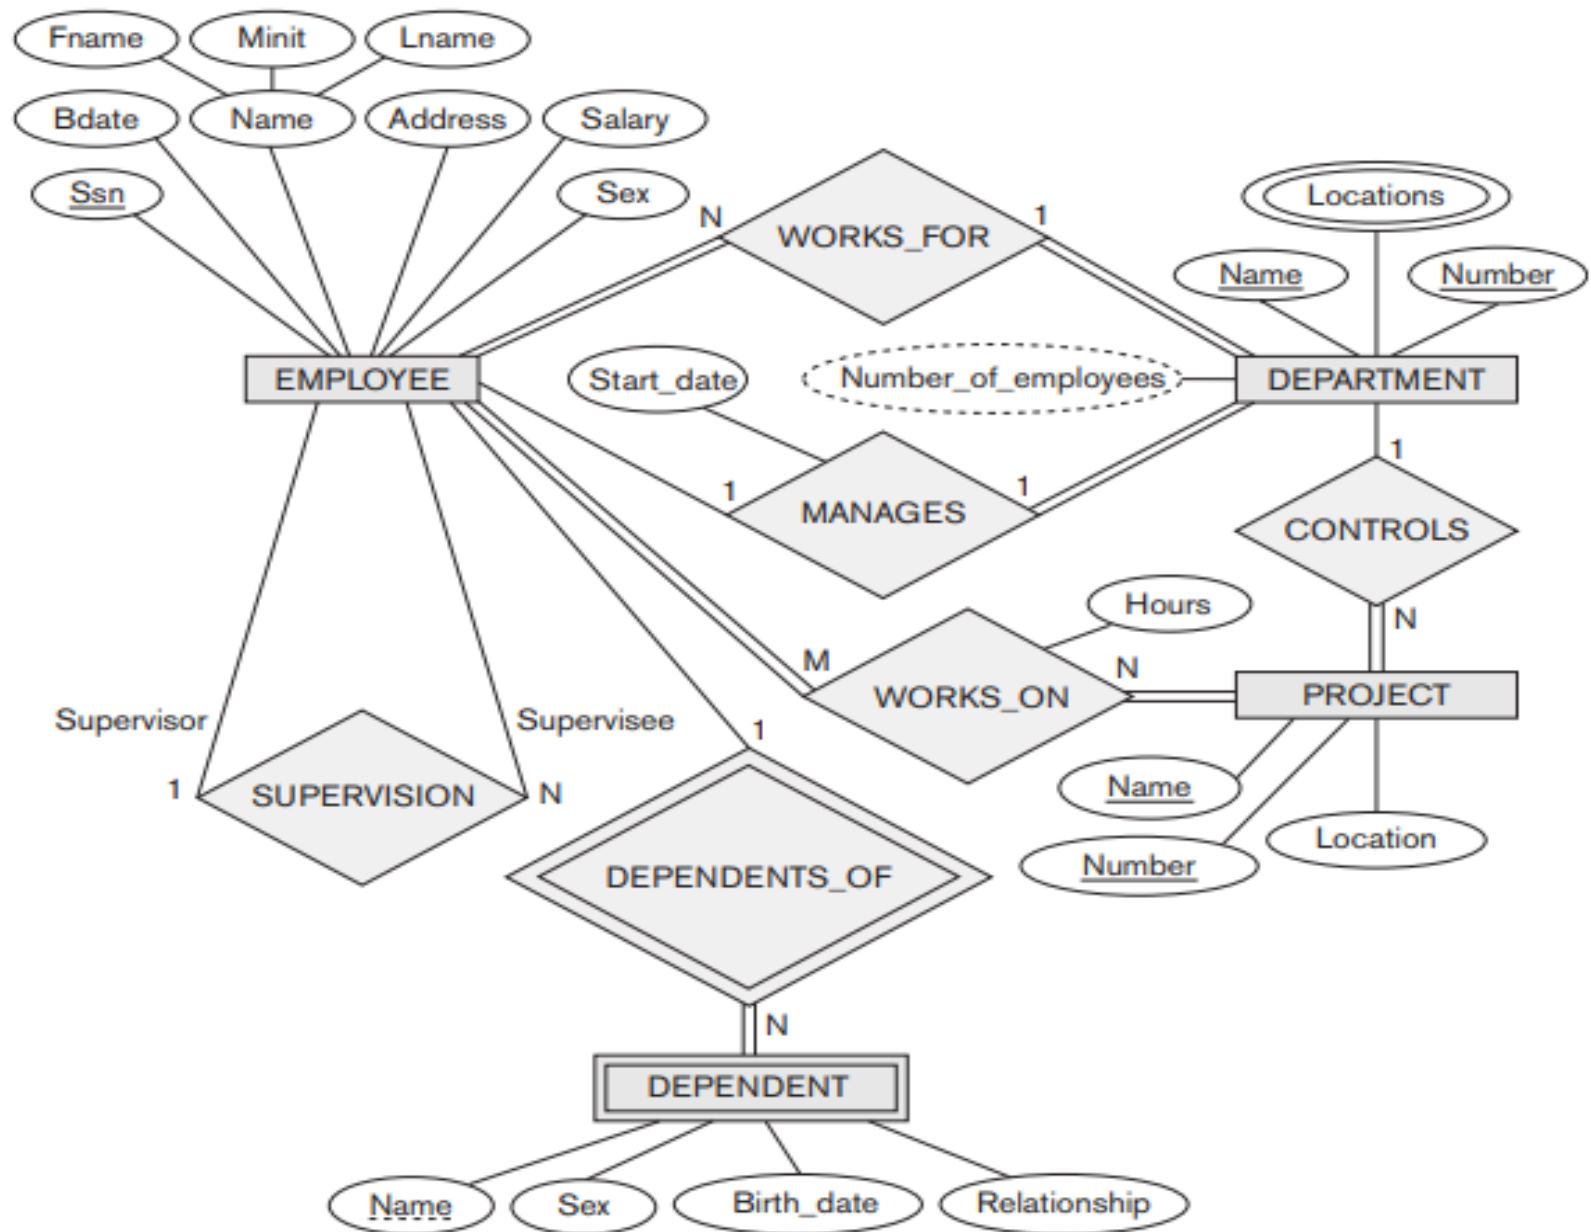
A Weak entity set

- has no key of its own;
- exists only because of association with strong entities

Example:




A complex ER Model



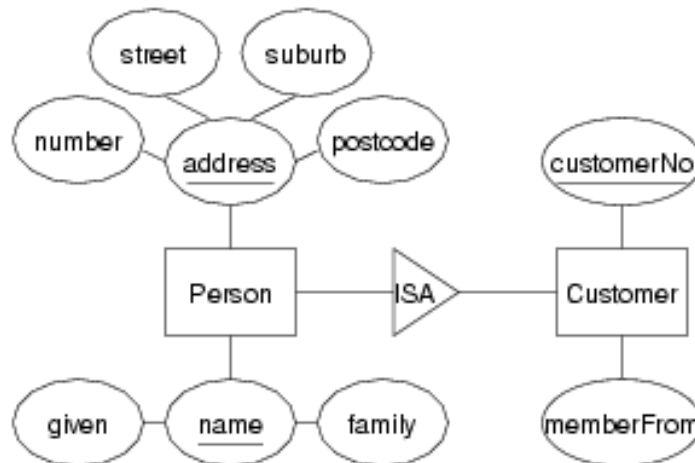
e.g. an ER model for a Bank

(1) Modelling people (employees)



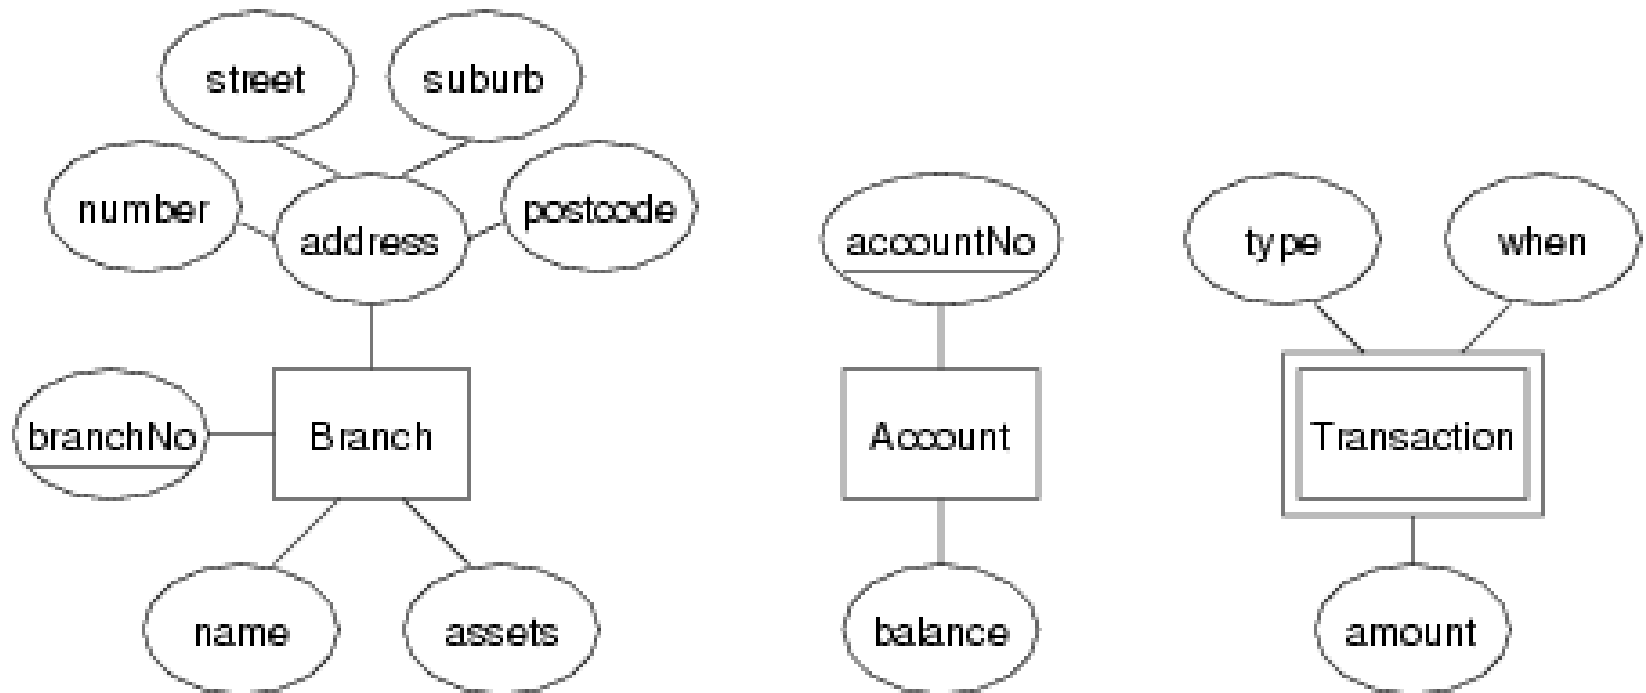
(2) Modelling people (customer)

Modelling people (cont):



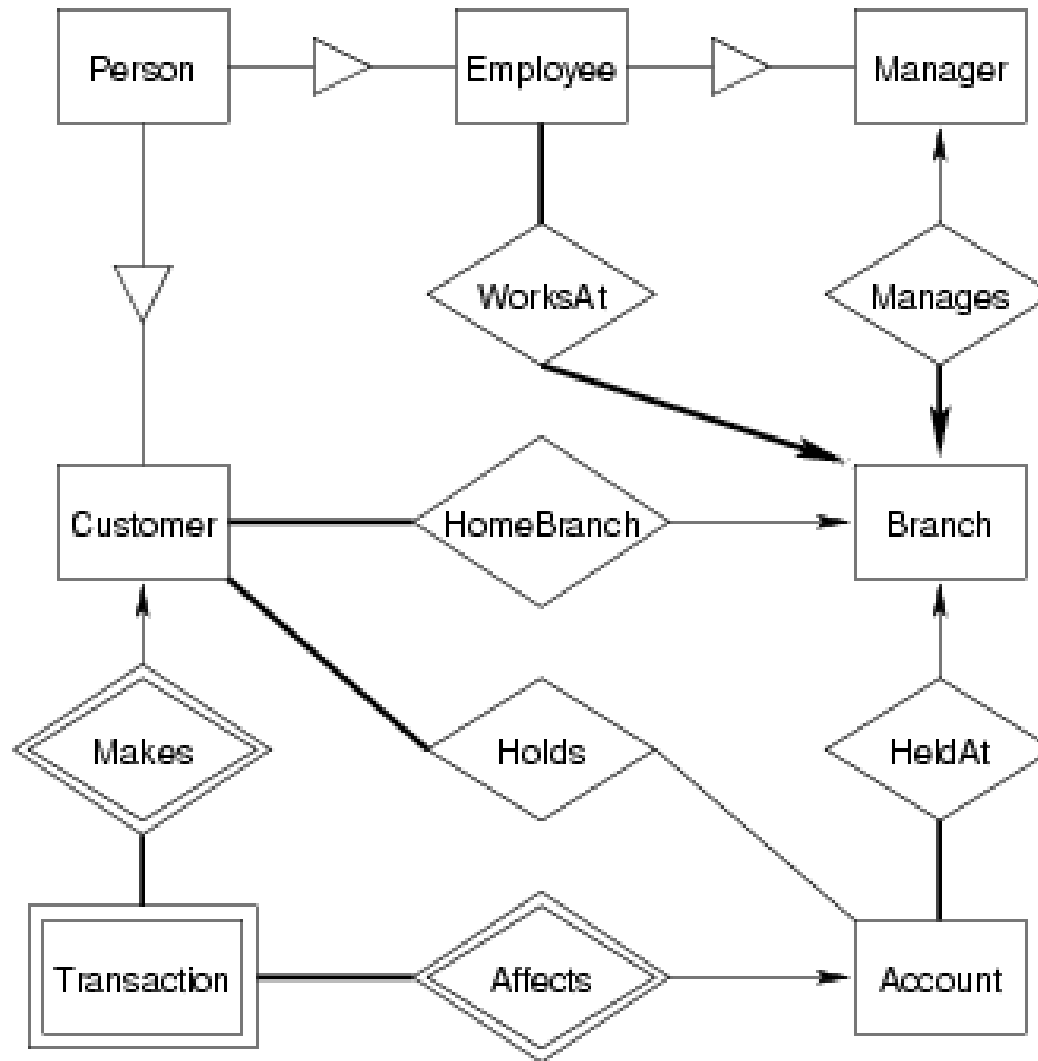
e.g. an ER model for a Bank

(3) Modelling branches, accounts, transactions



e.g. an ER model for a Bank

(4) Putting it all together with relationships



Exercise 2:

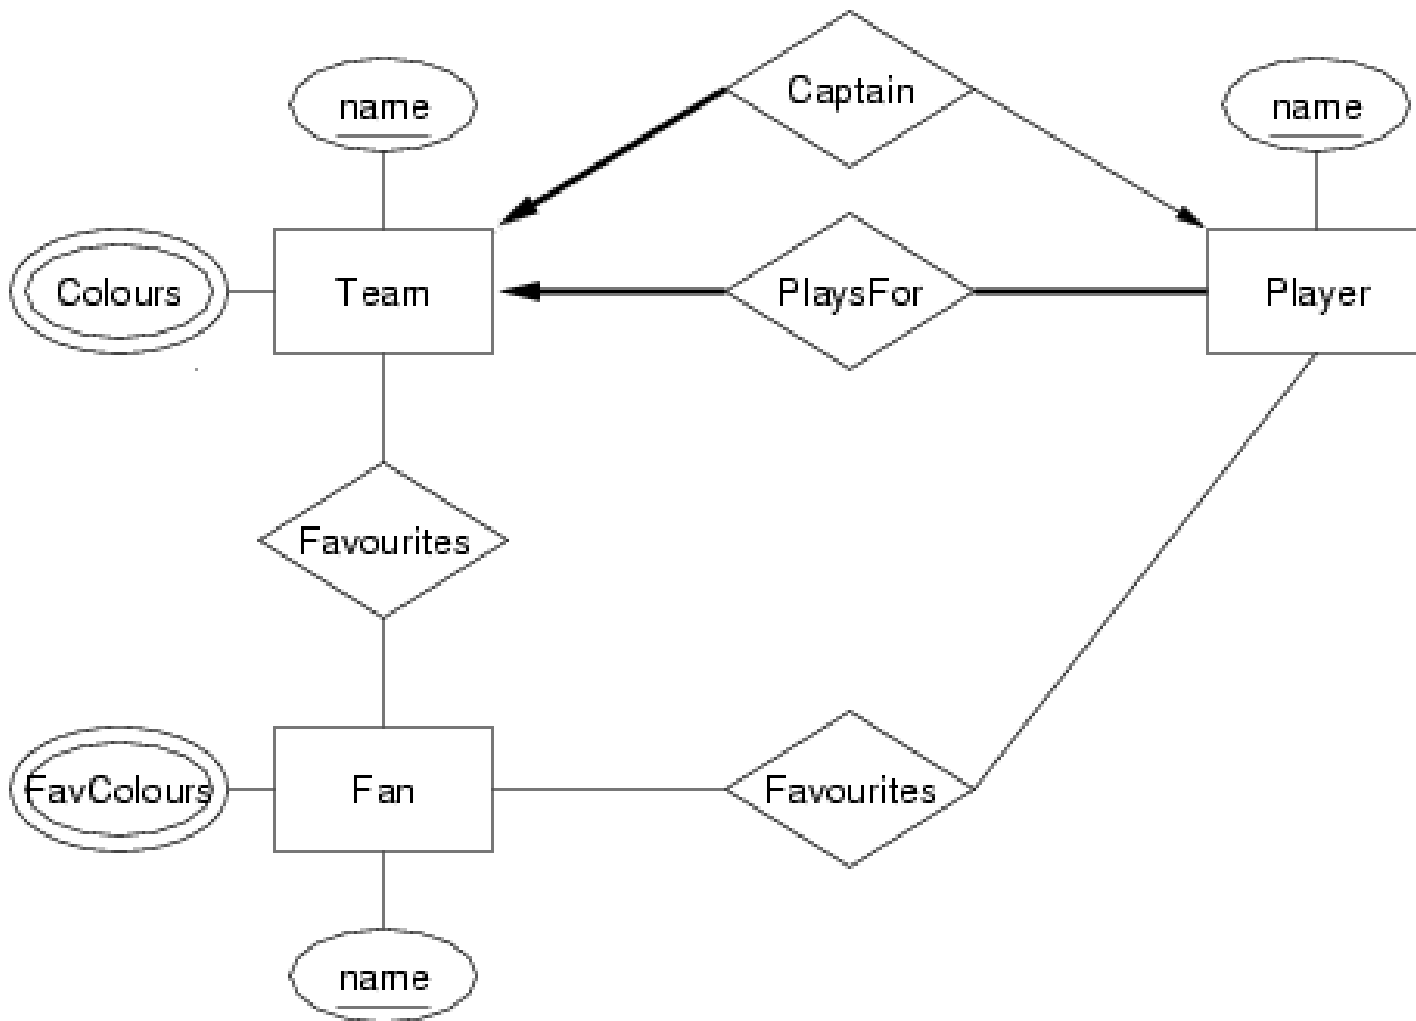
Develop an ER design for the following scenario:

A database records information about teams, players, and their fans, including:

- For each team, its name, its players, its captain (one of its players) and the colours of its uniform.
- For each player, their name and team.
- For each fan, their name, favourite teams, favourite players, and favourite colour.

State all assumptions made

Exercise 2 (Solution): ER Design



Exercise 3: Give an ER design to model the following scenario ...

- Patients are identified by an **SSN**, and their names, addresses and ages must be recorded.
- Doctors are identified by an SSN. For each doctor, the name, specialty and years of experience must be recorded.
- Each pharmacy has a name, address and phone number. A pharmacy **must** have a manager.
- A pharmacist is identified by an SSN, he/she can only work for one pharmacy. For each pharmacist, the name, qualification must be recorded.
- For each drug, the trade name and formula must be recorded.
- Every patient has a primary physician. Every doctor has at least one patient.
- Each pharmacy **sells** several drugs, and has a price for each. A drug could be sold at several pharmacies, and the price could vary between pharmacies.
- Doctors prescribe drugs for patients. A doctor could prescribe one or more drugs for several patients, and a patient could obtain prescriptions from several doctors. Each prescription has a date and quantity associated with it.

What kind of data, relationships and constraints exist?

Possible data:

- people: doctors, patients, pharmacists
- pharmacies, drugs
- person's SSN, name, address(?)
- doctor: as for person + specialty
- pharmacist: as for person + qualification etc.

Possible relationships:

- doctors **treat** patients, patients **have primary** physician
- pharmacists **work in** pharmacies
- drugs **are sold in** pharmacies
- doctors **prescribe** drugs **for** patients etc.

Possible constraints:

- every person has **exactly one, unique** SSN
- pharmacist works in ≤ 1 pharmacy
- patient has **exactly one** primary physician
- doctor treats ≥ 1 patient etc.

Summary of ER

- ER model is popular for doing conceptual design
 - high-level, models relatively easy to understand
 - good expressive power, can capture many details
- Basic constructs:
 - entities, relationships, attributes
 - relationship constraints: total / partial, $n:m$ / $1:n$ / $1:1$
- Other constructs:
 - inheritance hierarchies, weak entities
- Many notational variants of ER exist
(especially in the expression of constraints on relationships)

Relational Data Model

Relational Data Model

The **relational data model** describes the world as a collection of inter-connected **relations** (or **tables**)

Goal of relational model:

- a simple, general data modelling formalism
- maps easily to file structures (i.e. implementable)

Relational model has two styles of terminology:

mathematical	relation	tuple	attribute
data-oriented	table	record (row)	field (column)

STUDENT

Name	Student_number	Class	Major
Smith	17	1	CS
Brown	8	2	CS

Relational Data Model

The relational model has one structuring mechanism ...

- a relation corresponds to a mathematical "**relation**"
- a relation can also be viewed as a "**table**"

Each **relation** (table) (denoted R, S, T, \dots) has:

- a **name** (unique within a given database)
- a set of **attributes** (or column headings)

Each attribute (denoted A, B, \dots or a_1, a_2, \dots) has:

- a **name** (unique within a given relation)
- an associated **domain** (set of allowed values)

DB definitions also make extensive use of **constraints**

Relational Data Model

Example of a relation (table): Bank Account

The diagram illustrates a relation (table) named 'Account'. The table has three columns: 'branchName', 'accountNo', and 'balance'. The data rows are as follows:

branchName	accountNo	balance
Downtown	A-101	500
Mianus	A-215	700
Perryridge	A-102	400
Round Hill	A-305	350
Brighton	A-201	900
Redwood	A-222	700
Brighton	A-217	750

Annotations:

- An arrow points from the text *Relation, Table* to the word **Account**.
- Three arrows point from the text *Attributes, Columns, Fields* to the column headers: **branchName**, **accountNo**, and **balance**.
- Seven arrows point from the text *Tuples, Rows, Records* to each of the seven data rows in the table.

Relational Data Model

A **tuple (row)** is a **set** of **values (attribute or column values)**

Attribute values:

- Are **atomic** (no composite or multi-valued attributes).
- Belong to a **domain**; a domain has a **name**, **data type** and **format**
- A distinguished value **NULL** belongs to all domains.
- **NULL** has several interpretations: none, don't know, irrelevant

Column Header	Domain Name	Domain Data Type, Format and Constrain
phone_number	local_phone_numbers - (set of phone numbers valid in australia)	character string of the format <i>(dd) dddddddd</i> , where each <i>d</i> is a numeric (decimal) digit and the first two digits form a valid telephone area code.
age	employee_age (set of possible ages for employees in the company)	An integer value between 15 and 80

Relational Data Model

- A **relation(table)** is a **set** of **tuples**.
- Since a relation is a set, there is **no ordering** on rows.
- Normally, **we define a standard ordering** on components of a tuple.
- The following are different presentations of the same relation:

branchName	accountNo	balance
Downtown	A-101	500
Mianus	A-215	700
Perryridge	A-102	400
Round Hill	A-305	350
Redwood	A-222	700

accountNo	branchName	balance
A-305	Round Hill	350
A-222	Redwood	700
A-215	Mianus	700
A-102	Perryridge	400
A-101	Downtown	500

- Each relation generally has a **primary key** (subset of attributes, unique over relation)

Relational Data Model

A *database* is a **set** of *relations(tables)*

Account

branchName	accountNo	balance
Downtown	A-101	500
Mianus	A-215	700
Perryridge	A-102	400
Round Hill	A-305	350
Brighton	A-201	900
Redwood	A-222	700

Branch

branchName	address	assets
Downtown	Brooklyn	9000000
Redwood	Palo Alto	2100000
Perryridge	Horseneck	1700000
Mianus	Horseneck	400000
Round Hill	Horseneck	8000000
North Town	Rye	3700000
Brighton	Brooklyn	7100000

Customer

name	address	customerNo	homeBranch
Smith	Rye	1234567	Mianus
Jones	Palo Alto	9876543	Redwood
Smith	Brooklyn	1313131	Downtown
Curry	Rye	1111111	Mianus

Depositor

account	customer
A-101	1313131
A-215	1111111
A-102	1313131
A-305	1234567
A-201	9876543
A-222	1111111
A-102	1234567

Expressing Relational Data Model Mathematically

Given a relation (table) R which has:

- n attributes a_1, a_2, \dots, a_n
- with corresponding domains D_1, D_2, \dots, D_n

We define:

- **Relation Schema of R** as: $R(a_1:D_1, a_2:D_2, \dots, a_n:D_n)$
- **Tuple of R** as: an element of $D_1 \times D_2 \times \dots \times D_n$ (i.e. list of values)
- **Instance of R** as: subset of $D_1 \times D_2 \times \dots \times D_n$ (i.e. set of tuples)
- **Database schema** : a collection of relation schemas.
- **Database (instance)** : a collection of relation instances.

Example of a Relation Schema

Given a relation or table, **Account** , which has:

- 3 attributes `branchName`, `accountNo`, `balance`
- with corresponding domains `string`, `string`, `int` ,

then we can define the schema of **Account** as:

Account (`branchName:string`, `accountNo:string`, `balance:int`) OR simply as

Account (`branchName`, `accountNo`, `balance`)

and a tuple **Account** (i.e., row of **Account**)can be specified as:

(`Downtown`, `A-101`, `500`)

and an instance of relation **Account** (a set of tuples or rows) as:

* **No duplicates**

{ (`Downtown`, `A-101`, `500`), (`Mianus`, `A-215`, `700`),
(`Perryridge`, `A-102`, `400`), (`Round Hill`, `A-305`, `350`),
(`Brighton`, `A-201`, `900`), (`Redwood`, `A-222`, `700`)}
37

Relation Schema

- The **degree (or arity)** of a relation is the number of attributes n of its relation schema, so a relation schema R of degree n is denoted by $R(a_1, a_2, \dots, a_n)$.

e.g. A relation of degree seven, which stores information about university students, would contain seven attributes describing each student. as follows:

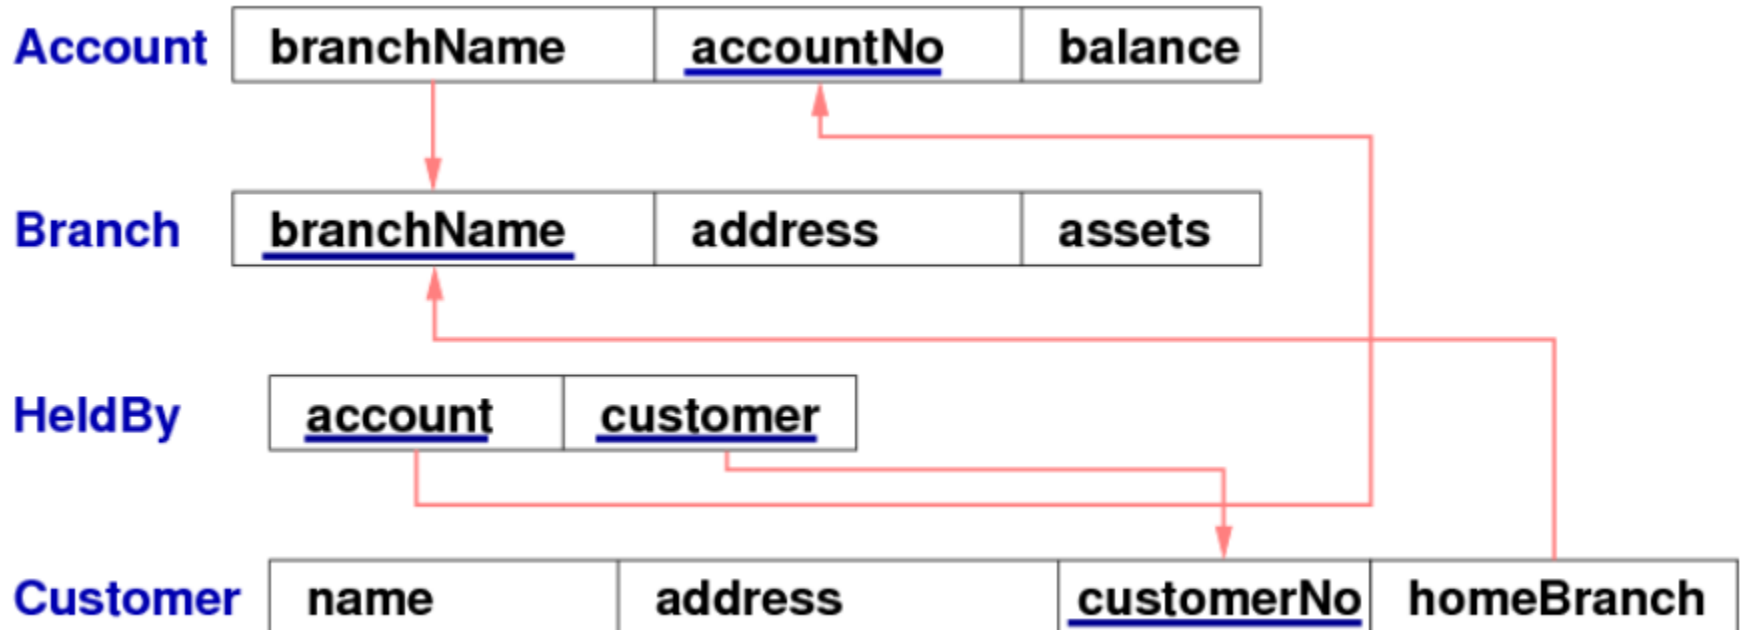
Student (name, ssn, home_phone, address, office_phone, age, gpa)

OR as

Student (name: string, ssn: string, home_phone: string, address: string, office_phone: string, age: integer, gpa: int)

Database Schema – a collection of relation schemas

Example of a **database schema** with 4 relation schemas:



Database Instance – a collection of relation instances

Example of a **database instance** with 4 relation instances:

Account

branchName	accountNo	balance
Downtown	A-101	500
Mianus	A-215	700
Perryridge	A-102	400
Round Hill	A-305	350
Brighton	A-201	900
Redwood	A-222	700

Branch

branchName	address	assets
Downtown	Brooklyn	9000000
Redwood	Palo Alto	2100000
Perryridge	Horseneck	1700000
Mianus	Horseneck	400000
Round Hill	Horseneck	8000000
North Town	Rye	3700000
Brighton	Brooklyn	7100000

Customer

name	address	customerNo	homeBranch
Smith	Rye	1234567	Mianus
Jones	Palo Alto	9876543	Redwood
Smith	Brooklyn	1313131	Downtown
Curry	Rye	1111111	Mianus

HeldBy

account	customer
A-101	1313131
A-215	1111111
A-102	1313131
A-305	1234567
A-201	9876543
A-222	1111111
A-102	1234567

Constraints

- **Relations** are used to represent real-world *entities* and *relationships* between these *entities*
- To represent real-world problems, need to describe
 - what **values are/are not** allowed
 - what **combinations of values are/are not** allowed
- **Constraints** are logical statements that do this:
 - Domain constraint
 - Key constraint
 - Referential integrity
- A DBMS should provide capabilities to enforce these constraints

Key (Uniqueness) Constraint

- A **key** attribute is one whose value can be used to uniquely identify each tuple (row) in the relation
- A relation can have more than one key, so each key is a **candidate key** e.g., **registration** and **VIN_Chassis_no**

Registration	VIN_Chassis_No	Make	Model
NSW XNCNC – 121	A695323	Mercedes	GLC Coupe
VIC ABC145 – 908	BX98765	Mercedes	C Class

Example of key constraints:

- Student(id, ...) is guaranteed unique
- Class(..., day, time, location,...) is unique

Entity Integrity Constraint

- One of the candidate keys is designated as the **primary key** of the relation
- An **entity integrity constraint** states that no primary key value can be NULL

Domain Constraints

Domain constraints specify that within each tuple, the value of each attribute (a_1, a_2, \dots, a_n) must be an atomic value from the corresponding domain D_1, D_2, \dots, D_n

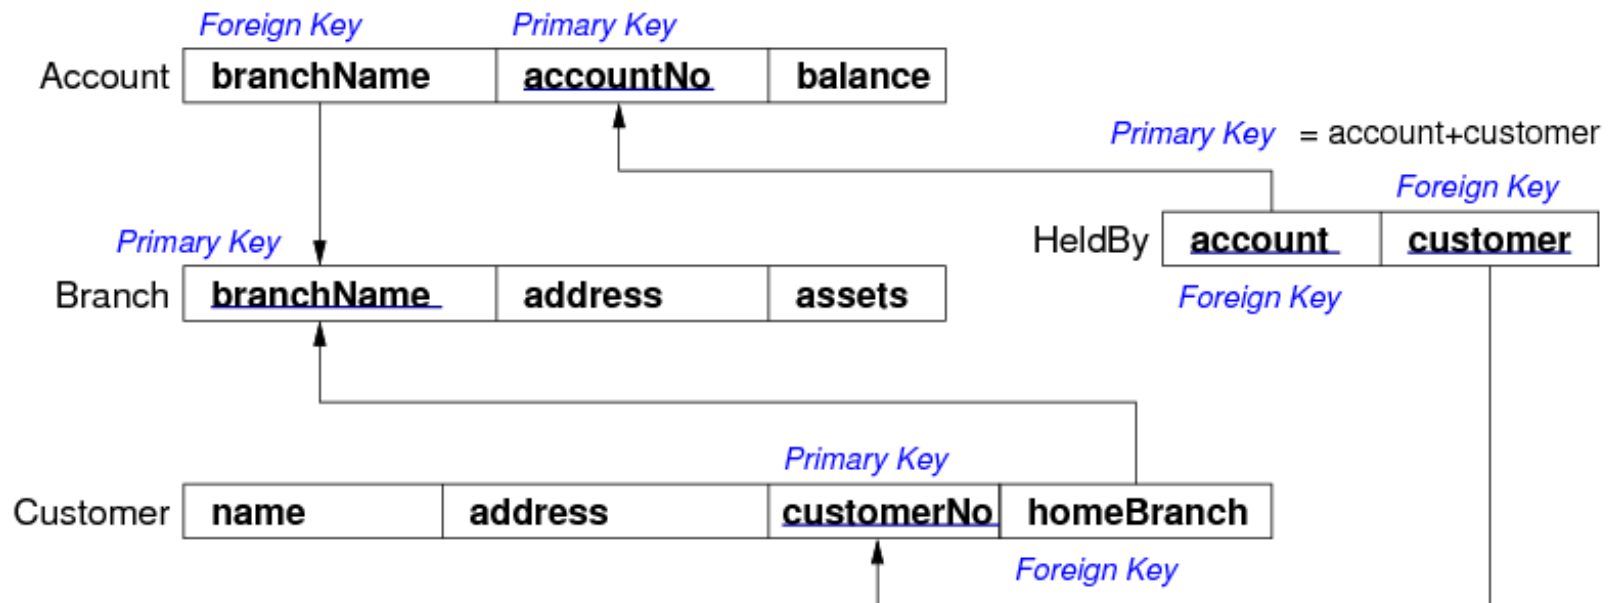
e.g.,

- **Employee.age** attribute is typically defined as **integer**
- better modelled by adding extra constraint **$(15 < \text{age} < 66)$**

Note: **NULL** satisfies all domain constraints (except (NOT NULL))

Referential Integrity Constraint

- describe references between relations (tables)
- are related to the notion of a foreign key (FK), where the **primary key** of the **parent** table is linked to a **foreign key** in the **child** table
e.g., the **Account** relation (**child**) needs to take note of the branch where each account is held
- the notion that the **branchName** (in Account) **must refer** to a valid **branchName** (in Branch) is a **referential integrity constraint**



Referential Integrity Constraints

- Related to the notion of a **foreign key**
- A set of attributes F in relation R_1 is a **foreign key** if:
 - the attributes in F **correspond** to the attributes in the primary key of another relation R_2
 - the value for F in each tuple of R_1
 - either occurs as a primary key in R_2
 - or is entirely NULL
- Foreign keys are critical in relational DBs;
 - they provide ...the "glue" that links individual relations (tables)
 - the way to assemble query answers from multiple tables

Putting all together...

A **relational database schema** is viewed as:

- a set of relation schemas $\{ R_1, R_2, \dots, R_n \}$, and
- a set of integrity constraints

A **relational database instance** is

- a set of relation instances $\{ r_1(R_1), r_2(R_2), \dots, r_n(R_n) \}$
- where all of the integrity constraints are satisfied

One of the important functions of a relational DBMS:

- ensure that all data in the database satisfies constraints

Relational Model vs DBMS

The relational model is a **mathematical construct**

- giving a representation for data structures
- with constraints on relations/tuples
- and an *algebra* for manipulating relations/tuples (union, intersect...)

Relational Database Management Systems (RDBMS)

- provide an **implementation** of the relational model
- Uses SQL (Structured Query Language) as language for data definition, query, updates

Describing Relational Schemas

- SQL (**Structured Query Language**) provides the formalism to express relational schemas
- SQL provides a **Data Definition Language (DDL)** for creating relations

```
CREATE TABLE TableName (  
  attrName1 domain1 constraints1 ,  
  attrName2 domain2 constraints2 , ...  
  PRIMARY KEY (attri,attrj,...)   
  FOREIGN KEY (attrx,attry,...)   
  REFERENCES OtherTable (attrm,attrn,...)   
);
```


DBMS Terminology

To remember:

- DBMS-level ... database names must be unique
- database-level ... schema names must be unique
- schema-level ... table names must be unique
- table-level ... attribute names must be unique

Sometimes it is convenient to use same name in several tables

We distinguish which attribute we mean using qualified names

e.g. **Account.branchName** vs **Branch.branchName**

Features of RDBMS

- Support large-scale data-intensive applications
- Provide efficient storage and retrieval (disk/memory management)
- Support multiple simultaneous users (privilege, protection)
- Support multiple simultaneous operations (transactions, concurrency)
- Maintain reliable access to the stored data (backup, recovery)
- Use **SQL** as language for:
 - data definition (creating, deleting relations i.e. tables)
 - relation query (selecting tuples)
 - relation update (changing relations)

Mapping ER to Relational Data Model

Tomorrow ...

Mapping ER Designs to Relational Model

ER to Relational Mapping

By examining semantic correspondences, a formal mapping between the ER and relational models has been developed

ER Model	Relational Model
ER attribute	attribute (atomic)
ER entity-instance ER relationship-instance	tuple (row)
ER entity-set ER relationship	relation (table)
ER key	primary key of relation

Relational Model vs Entity Model

There are also differences between relational and ER models
Compared to ER, the relational model:

- Uses *relations* to model *entities* and *relationships*
- Has **no** *composite* or *multi-valued* attributes (only atomic)
- Has **no** *object-oriented notions* (e.g. subclasses, inheritance)

(1) Mapping Strong Entities

An *entity* consists of:

- a collection of attributes;
attributes are simple, composite and multi-valued

A *relation schema* consists of:

- a collection of attributes;
all attributes have **atomic** data values

So, even the **mapping** from entity to relation schema is **not simple**


(1) Mapping Strong Entities

An obvious mapping

- an entity set E with **atomic attributes** a_1, a_2, \dots, A_n **maps to**
- a relation (table) R with **attributes (columns)** a_1, a_2, \dots, A_n

Each row in relation R corresponds to an entity in E

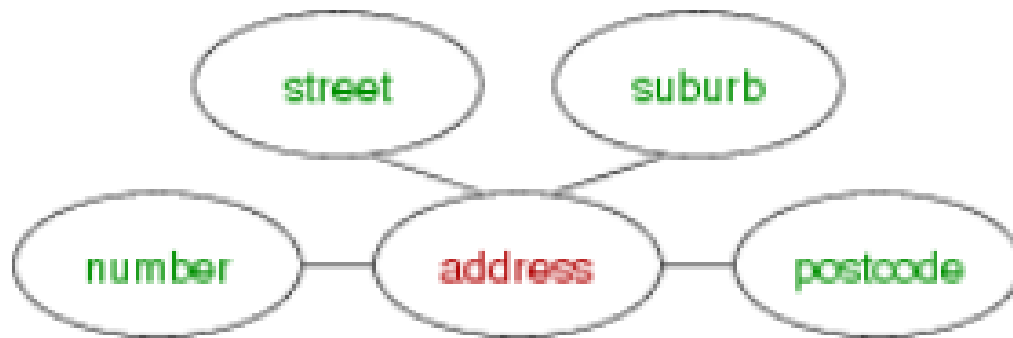
Example:



(Note: the key is preserved in the mapping)

(2) Mapping Composite Attributes

- ER supports composite (hierarchical) attributes.
- The relational model supports only atomic attributes.
- Composite attributes consist of
 - structuring attributes (non-leaf attributes)
 - data attributes (containing atomic values)



(2) Mapping Composite Attributes

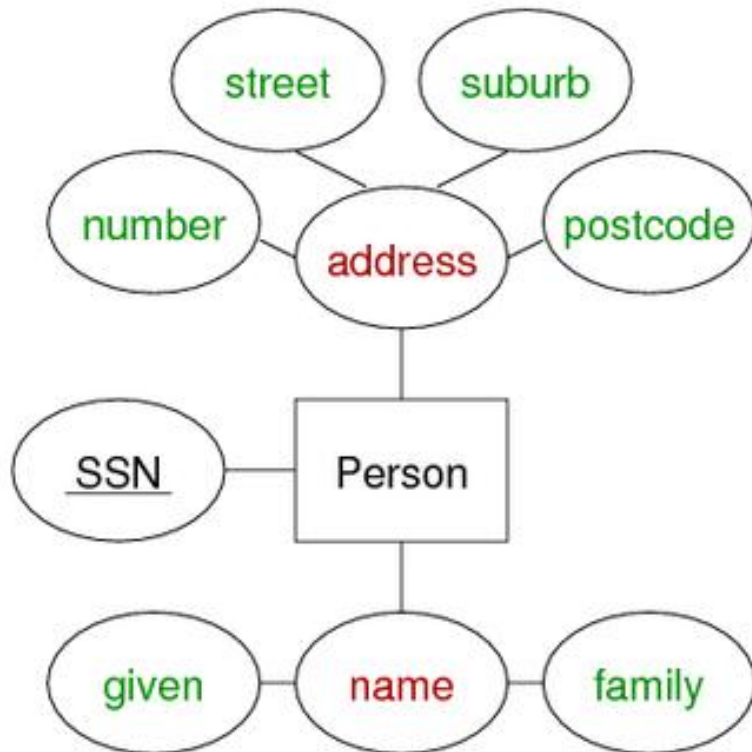
- One approach : remove structuring attributes
 - map atomic components to a set of atomic attributes (possibly with renaming)
 - E.g. Addr {number, street, suburb, pcode} maps to (AddrNumber, AddrStreet, AddrSuburb, AddrPcode)
- Alternative approach: concatenate atomic attribute values into a string
 - e.g., name {"John","Smith"} → "John Smith"
 - However, this approach:
 - requires extra extraction effort if components *are* required
 - cannot exploit efficient query capabilities on components

(2) Mapping composite attributes

Mapped by concatenation or flattening

Example:

ER Model



Relational Version #1

Person

<u>SSN</u>	name	address
------------	------	---------

Relational Version #2

Person

<u>SSN</u>	givenName	familyName	...	
...	number	street	suburb	postcode

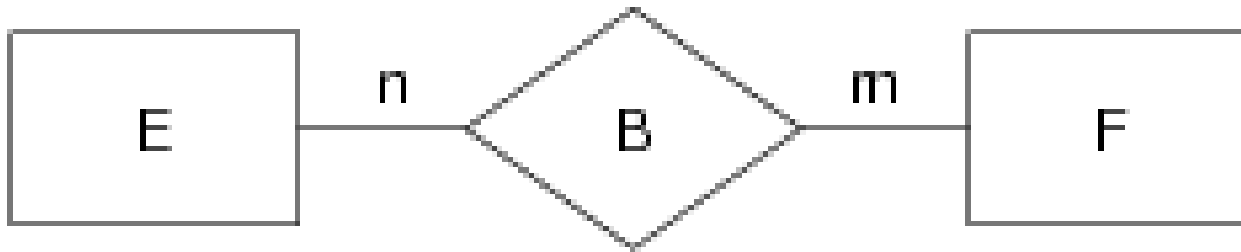
(3) Mapping Relationships

ER **relationship** → relational **table** (relation)

- Identify one entity as “parent” and other entity as “child”
- as general rule,
 - **PK of parent is added to child as FK**
- Any attributes of the relationship
 - are added to **child** relation

(3a) Mapping N:M Relationships

A binary relationship set B between entity sets E and F gives **associations** between pairs of entities in E and F



We can represent

- entity set E by relation S (using attribute mappings as above)
- entity set F by relation T (using attribute mappings as above)

But how to represent B ?

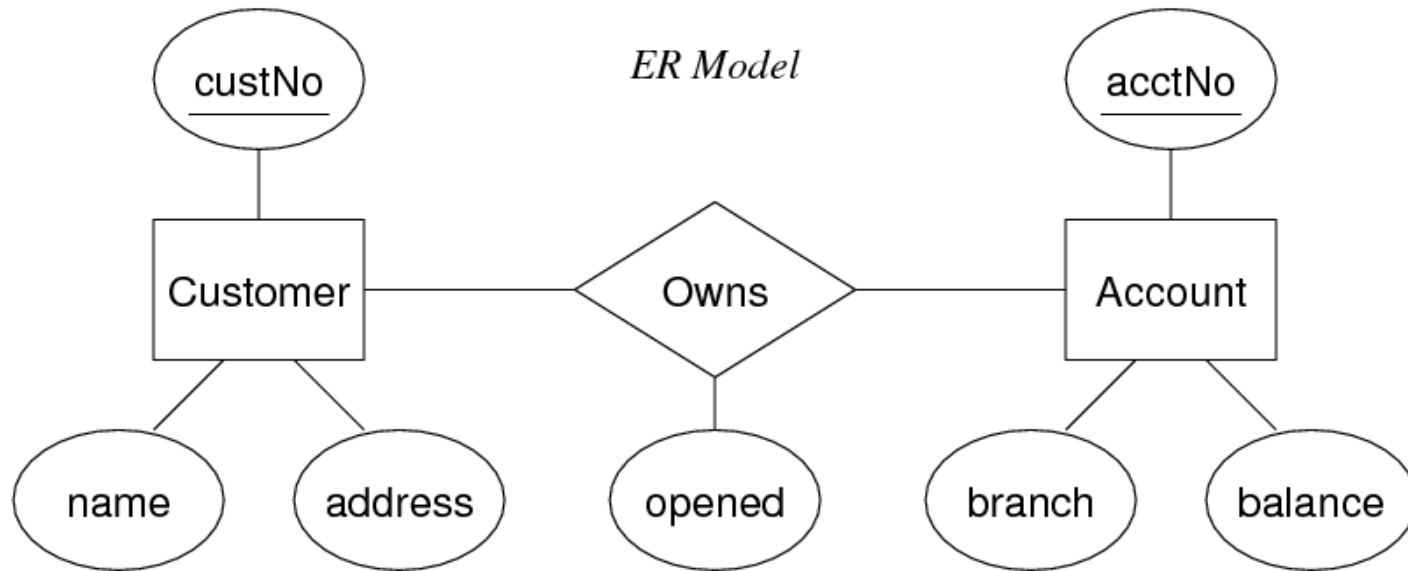
One possibility –

- Represent the **relationship** set B explicitly by a **relation** R containing:
 - all attributes from the primary keys of S and T
 - all attributes associated with the relationship set B
- where S and T are relations representing entity sets E and F
- and the **key** for R is the **union of the key attributes for S and T**

And this approach works generally for:

- relationship degree ≥ 2
- relationship multiplicity **1:1**, **1:N**, **N:M**
- associated attributes are simply included in R

Example - Mapping N:M Relationship



Relational Version

Customer

<u>custNo</u>	name	address
---------------	------	---------

Account

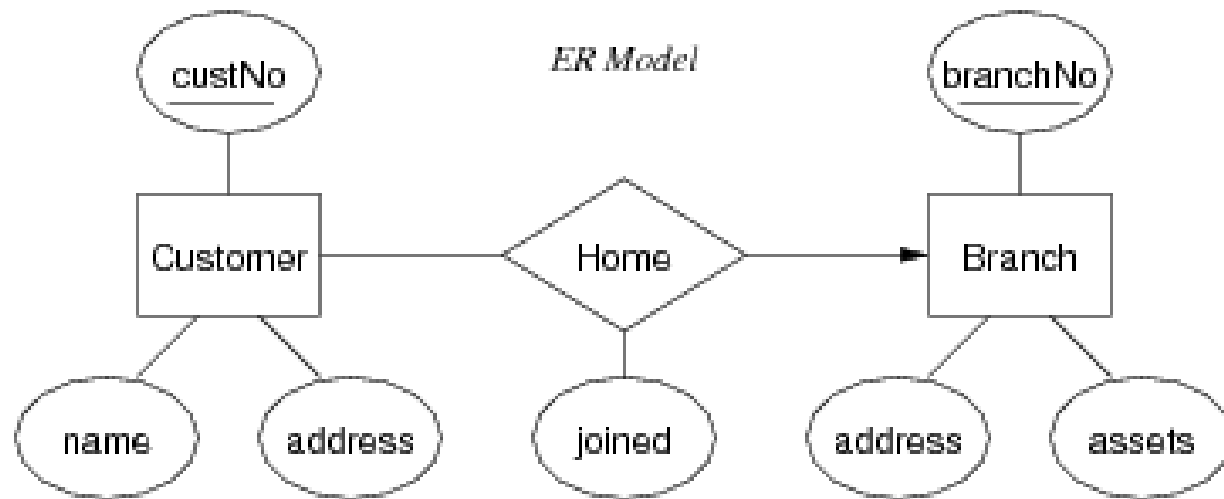
<u>acctNo</u>	branch	balance
---------------	--------	---------

Owns

<u>custNo</u>	<u>acctNo</u>	opened
---------------	---------------	--------

(3b) Mapping 1:M Relationships

Example:



Relational Version

Generic Mapping

Customer

<u>custNo</u>	name	address
---------------	------	---------

Branch

<u>branchNo</u>	address	assets
-----------------	---------	--------

Home

<u>custNo</u>	<u>branchNo</u>	joined
---------------	-----------------	--------

Optimised Mapping

Customer

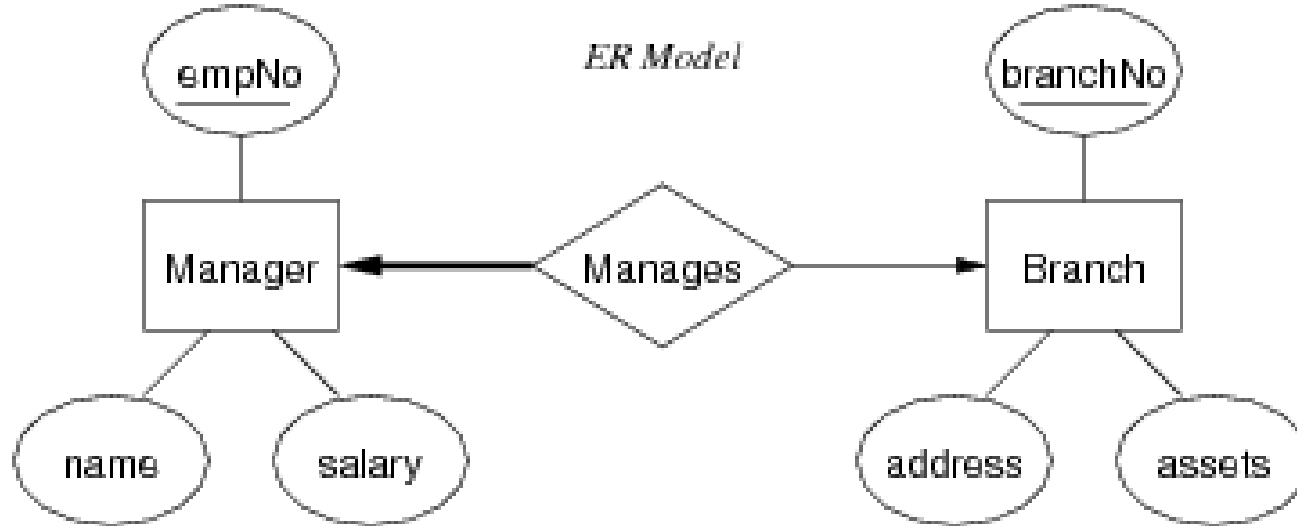
<u>custNo</u>	name	address	branchNo	joined
---------------	------	---------	----------	--------

Branch

<u>branchNo</u>	address	assets
-----------------	---------	--------

(3c) Mapping 1:1 Relationships

Example:

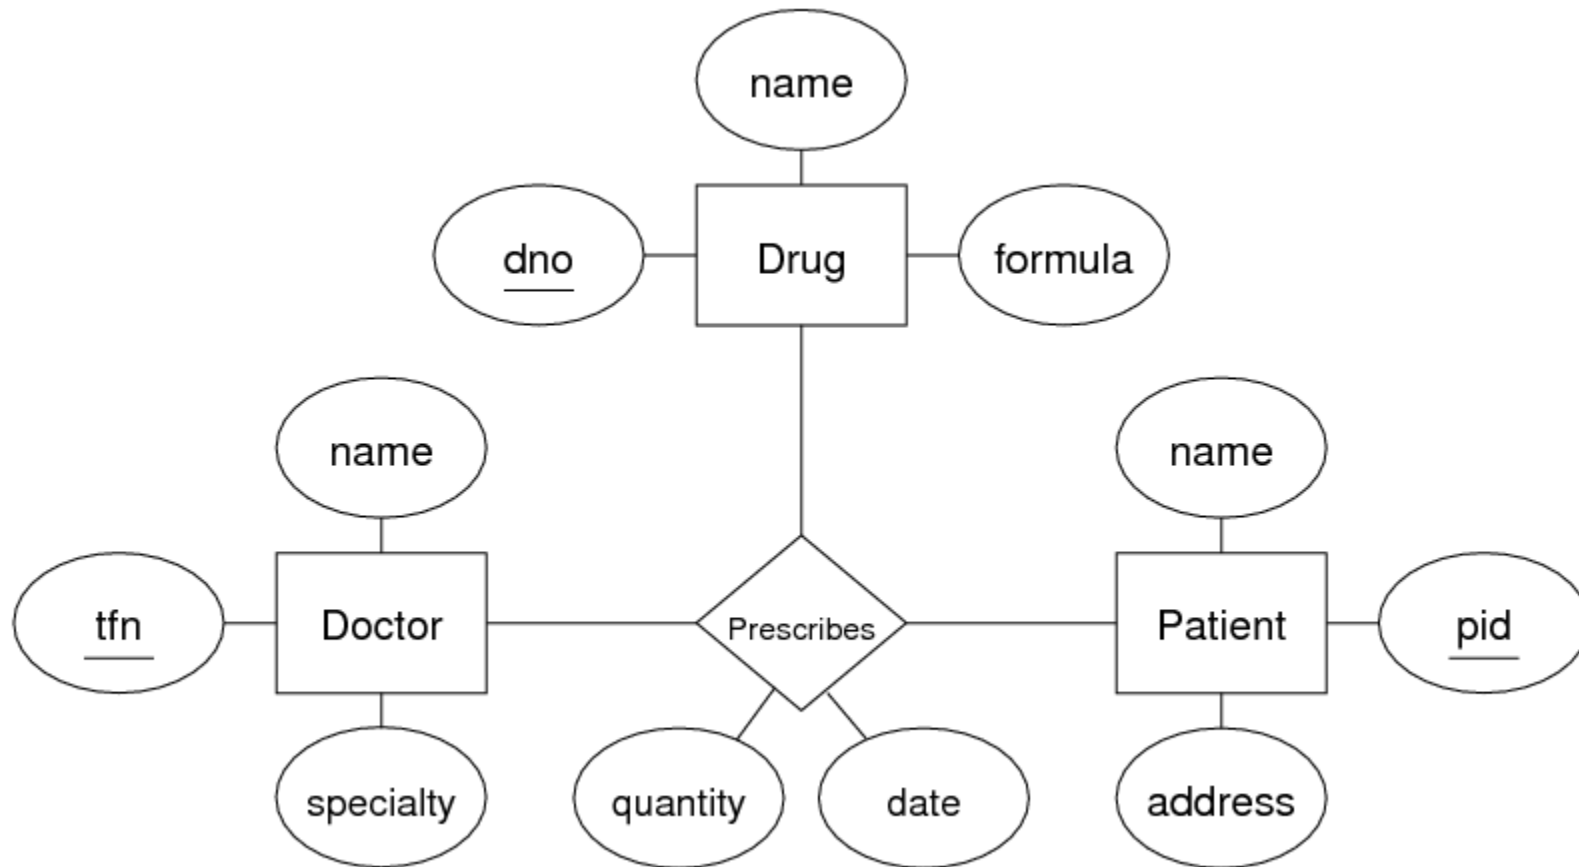


- Handled similarly to 1:N relationships
- For a 1:1 relationship between entity sets E and F (S and T):
 - choose one of S and T (e.g. S) (*Note : Choose the entity set that participates totally, if only one of them does*)
 - add the attributes of T 's primary key to S as foreign key
 - add the relationship attributes as attributes of S

(3d) Mapping n-way relationships

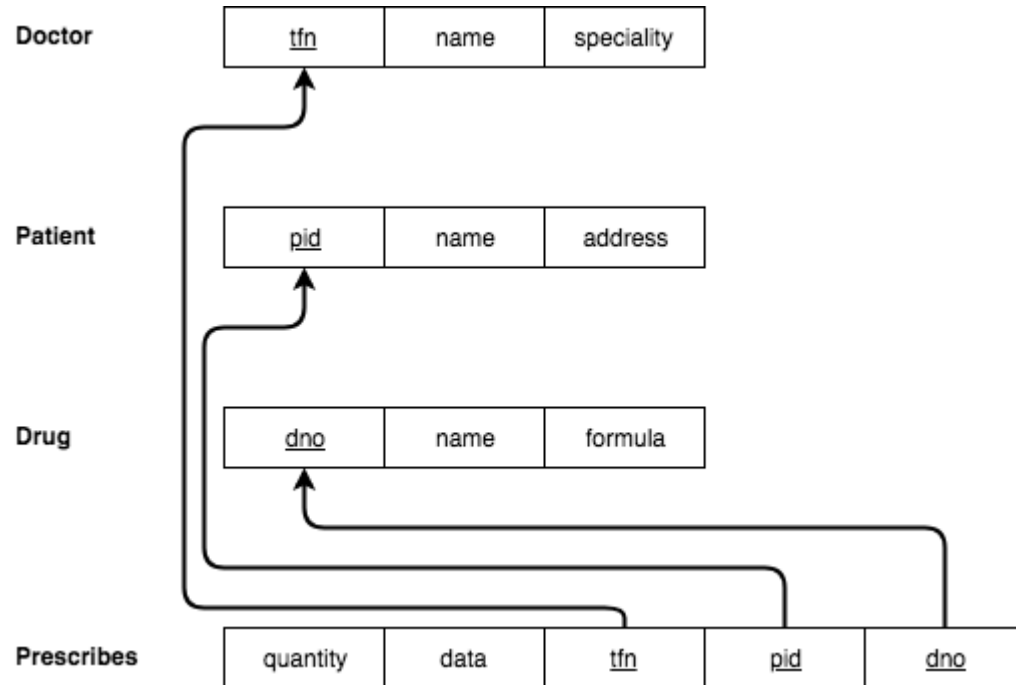
Exercise:

Convert the following ER design into a relational data model



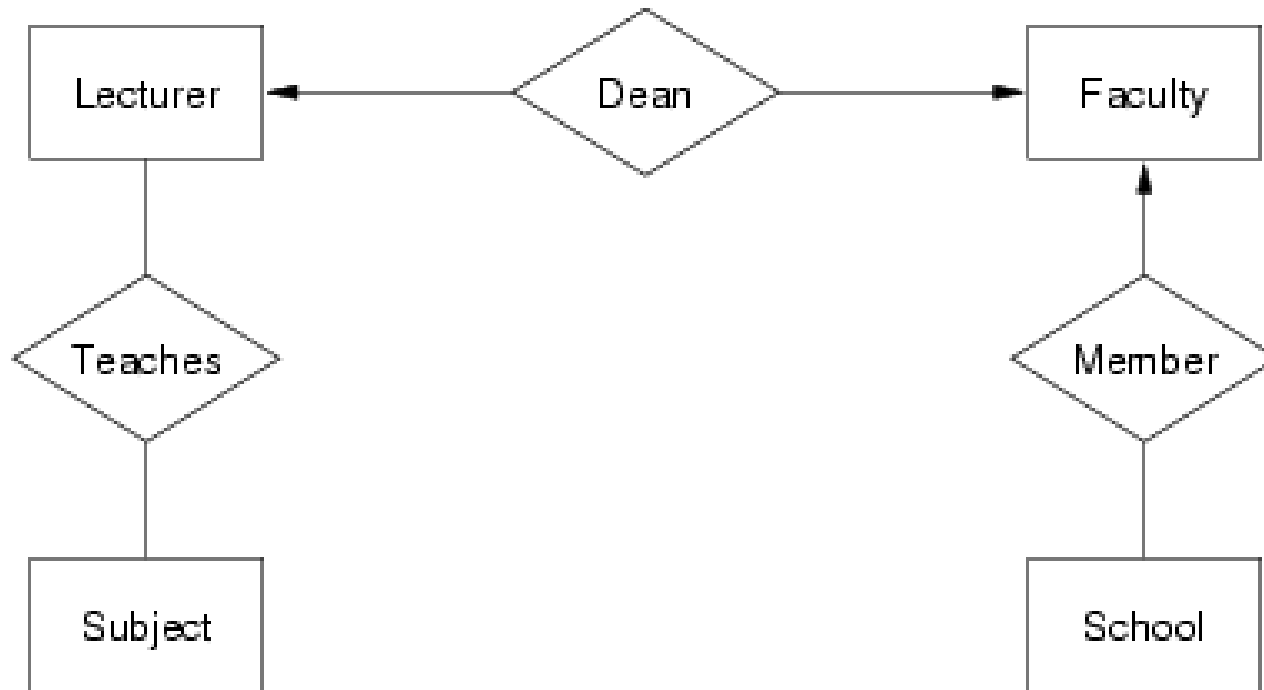
(3d) Mapping n-way relationships

Solution:




Exercise:

Convert the following ER design into a relational data model



You can assume that each attribute contains (at least) a suitably-named attribute containing a unique identifying number (e.g. the Lecturer entity contains a LecID attribute).

Solution: Relational Model



Relational model for a very small University ER model

So far ...

Summarising ER → Relational Mapping

Mapping entities and attributes

- ER **attribute** → relational **attribute**
- ER **entity** → relational **tuple**
- ER **entity-set** → relational **table** (relation)
- ER **key** → relational **primary key**

Mapping Relationships

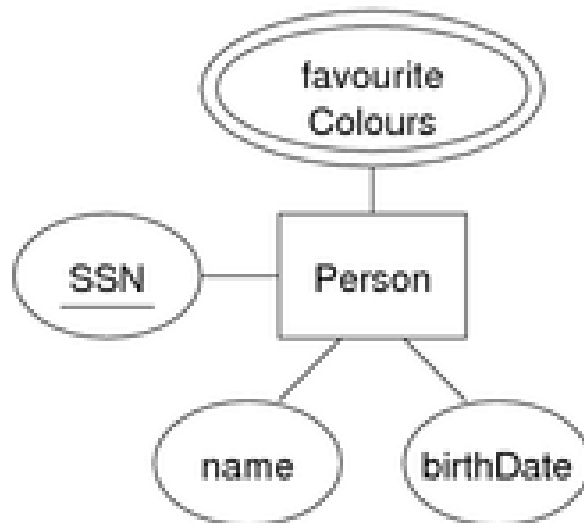
- ER **relationship** → create a **new** relational **table** (relation)
 - N:M relationship → add FK for each participating entity plus relationship attributes)
 - 1:N relationship → FK plus relationship attributes
 - 1:1 relationship → FK plus relationship attributes

(4) Mapping multi-valued attributes

- treat like an **N:M relationship** between **entities** and **values**
- create a **new relation** where each tuple contains:
 - the primary key attributes from the entity
 - one value for the multi-valued attribute from the corresponding entity

Example:

ER Model



Relational Version

Person

<u>SSN</u>	name	birthdate
------------	------	-----------

FavColour

<u>SSN</u>	<u>colour</u>
------------	---------------

(4) Mapping multi-valued attributes contd...

Example: the two entities

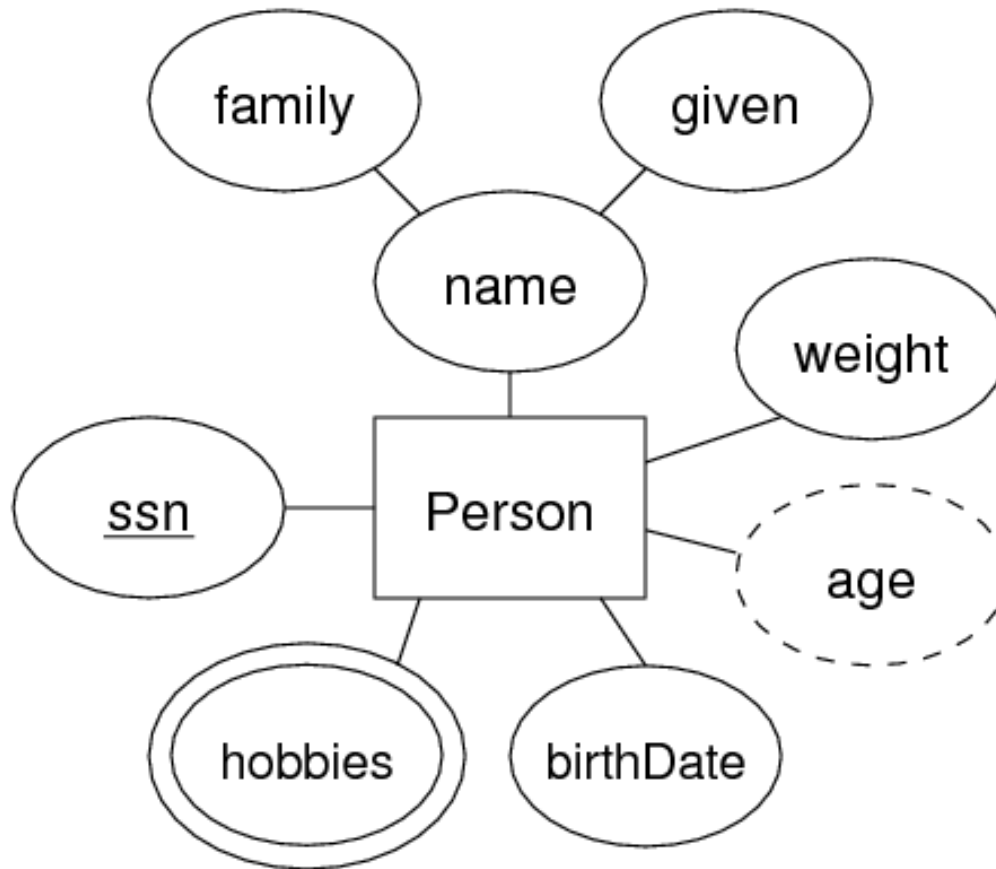
```
Person(12345, John, 12-feb-1990, [red, green, blue])  
Person(54321, Jane, 25-dec-1990, [green, purple])
```

would be represented as:

```
Person(12345, John, 12-feb-1990)  
Person(54321, Jane, 25-dec-1990)  
  
FavColour (12345, red)  
FavColour(12345, green)  
FavColour(12345, blue)  
FavColour(54321, green)  
FavColour(54321, purple)
```

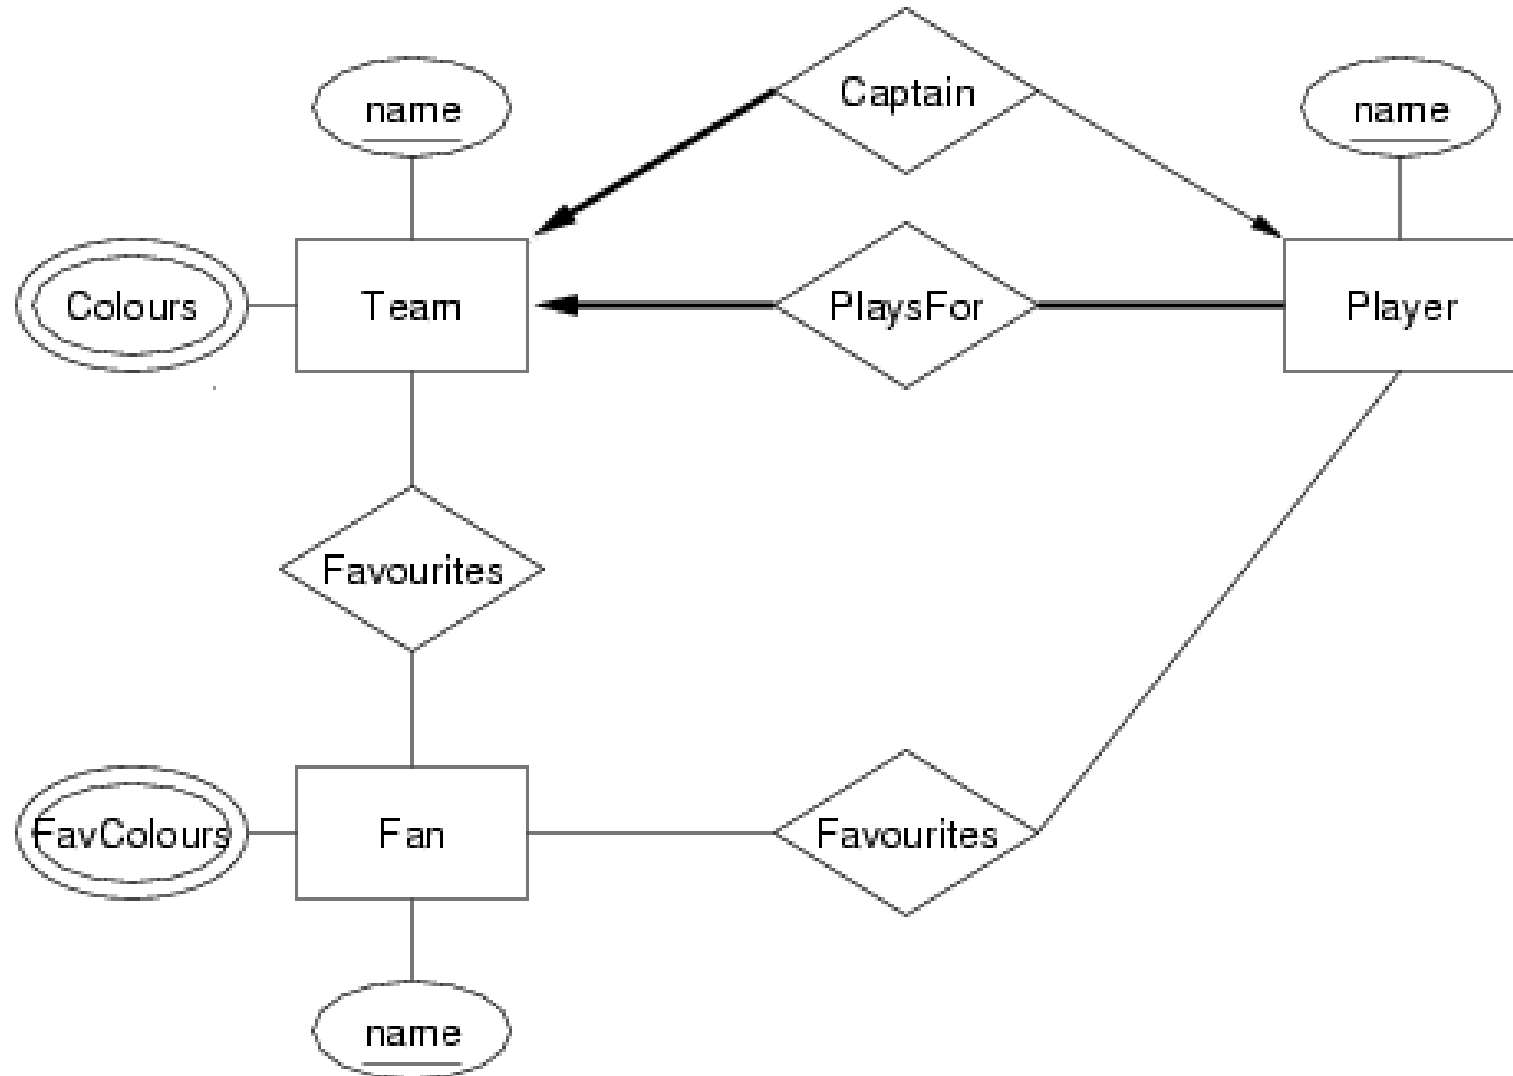

Exercise:

Convert the following ER design to relational form



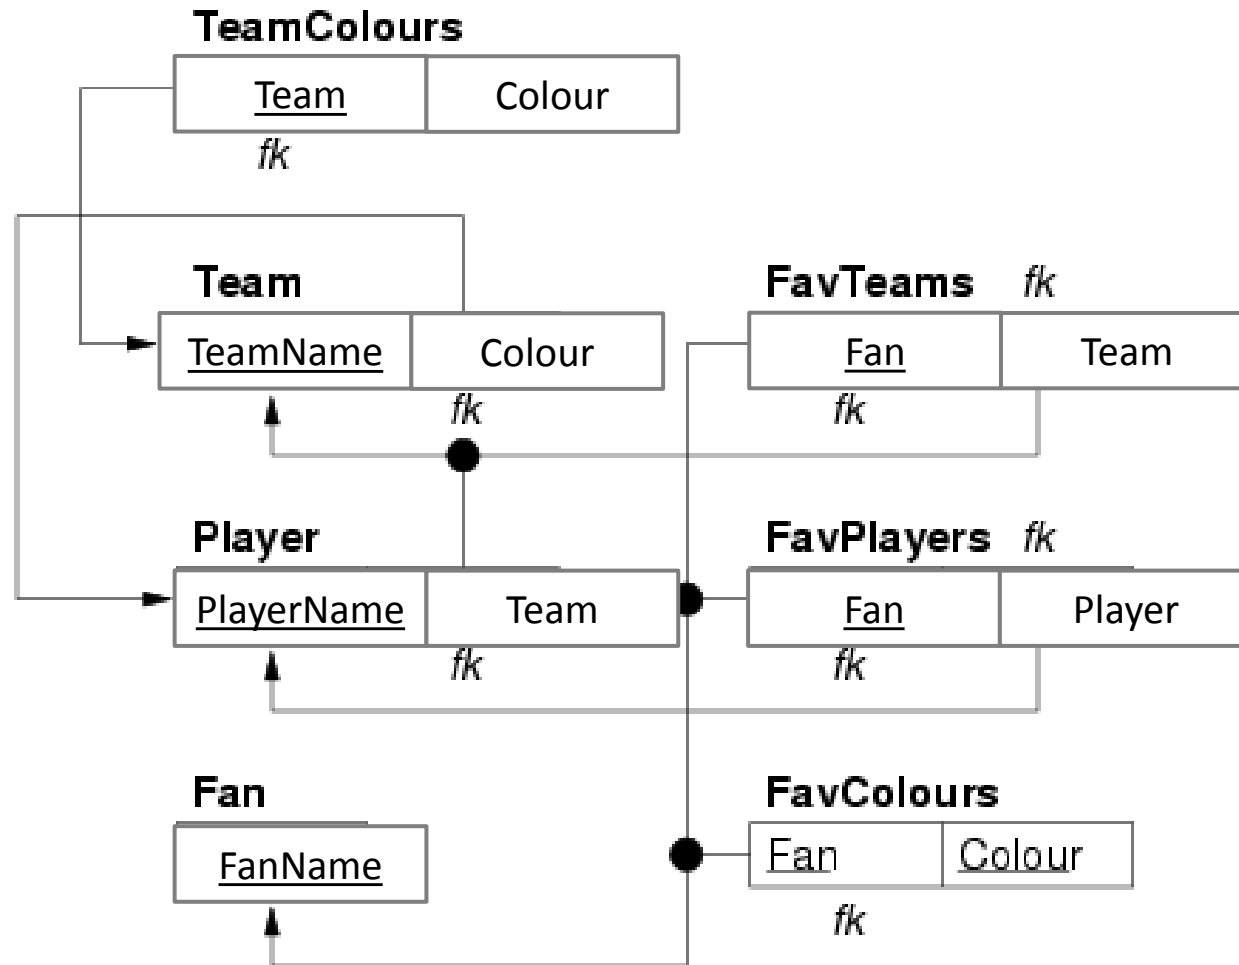
Exercise:

1. Convert the following ER design into a relational data model based on the box schema notation
2. Which elements of the ER design do not appear in the relational model?



Solution:

Convert the following ER design into a relational data model based on the box schema notation



(5) Mapping sub classes

Three different approaches to mapping subclasses to tables:


- **ER style**
 - each entity becomes a separate table,
 - containing attributes of subclass + FK to superclass table
- **object-oriented**
 - each entity becomes a separate table,
 - inheriting all attributes from all superclasses
- **single table with nulls**
 - whole class hierarchy becomes one table,
 - containing all attributes of all subclasses (null, if unused)

Which mapping is best depends on how data is to be used

(5) Mapping sub classes in ER style

The subclass relation contains:

- all of the subclass-specific attributes
- uses the superclass primary key to capture the association



Relational Model

Person

SSN	name	address
------------	-------------	----------------

Employee

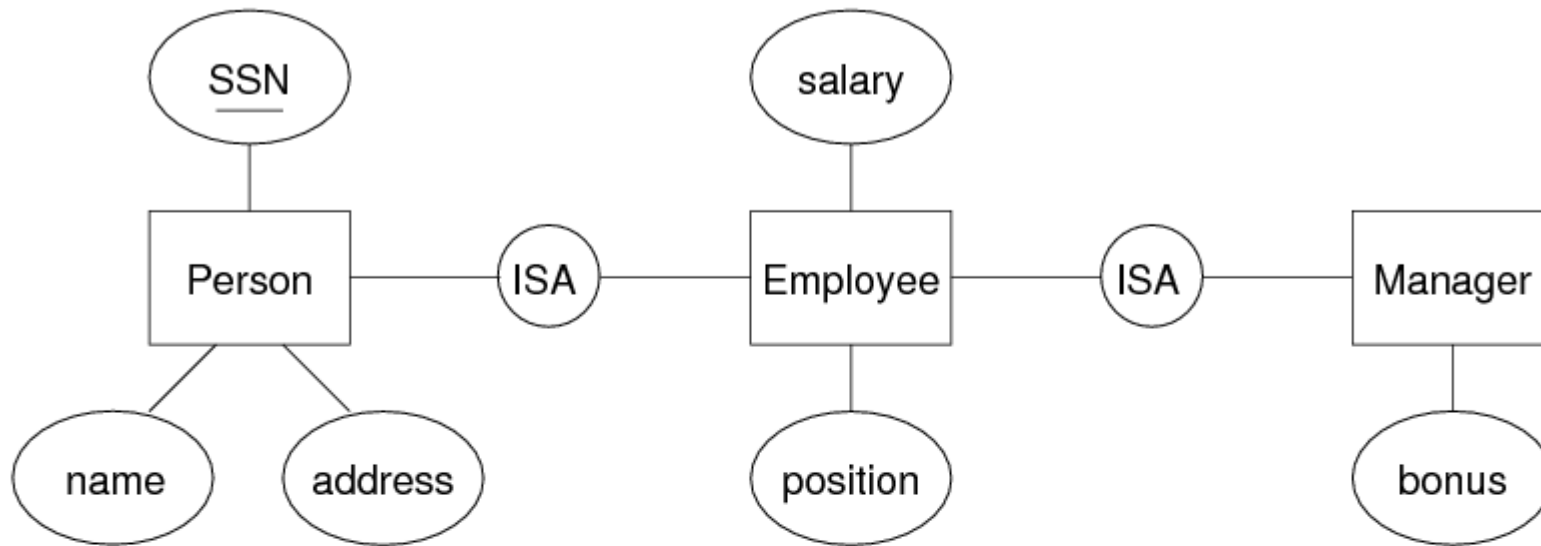
SSN	salary	position
------------	---------------	-----------------

Manager

SSN	bonus
------------	--------------

(5) OO Mapping of sub classes

Entity-Relationship Model



Relational Model

Person

<u>SSN</u>	name	address
------------	------	---------

Employee

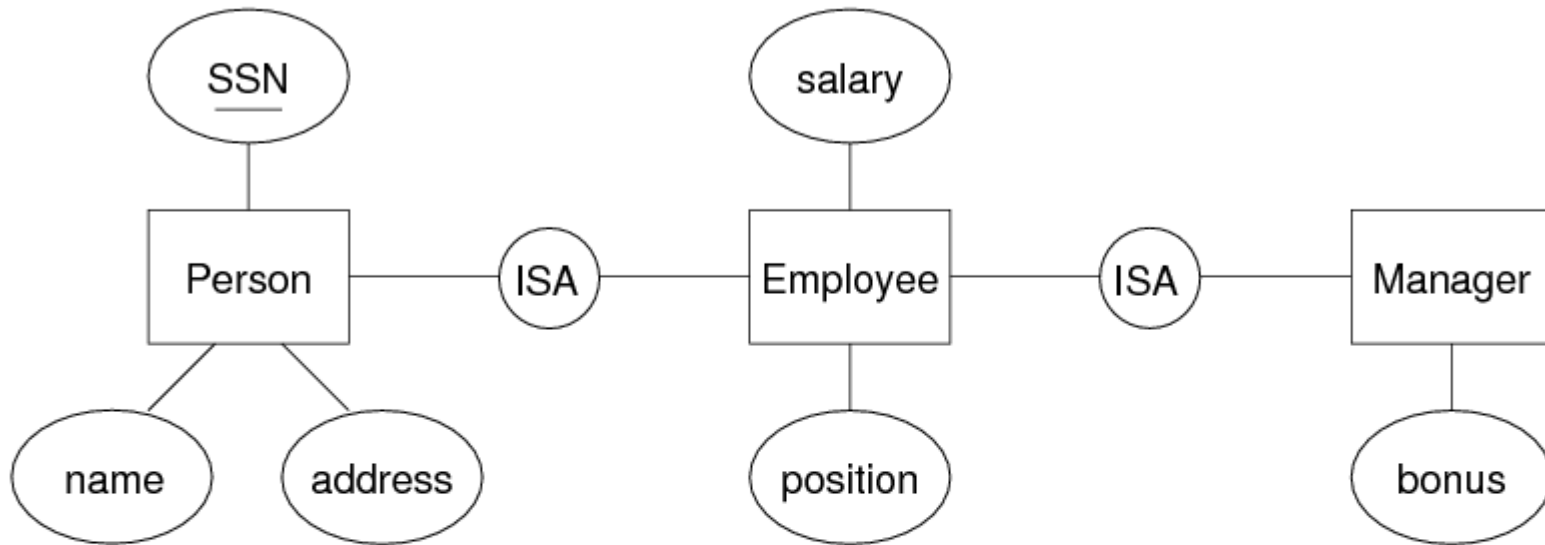
<u>SSN</u>	name	address	salary	position
------------	------	---------	--------	----------

Manager

<u>SSN</u>	name	address	salary	position	bonus
------------	------	---------	--------	----------	-------

(5) Single table with nulls mapping

Entity-Relationship Model



Relational Model

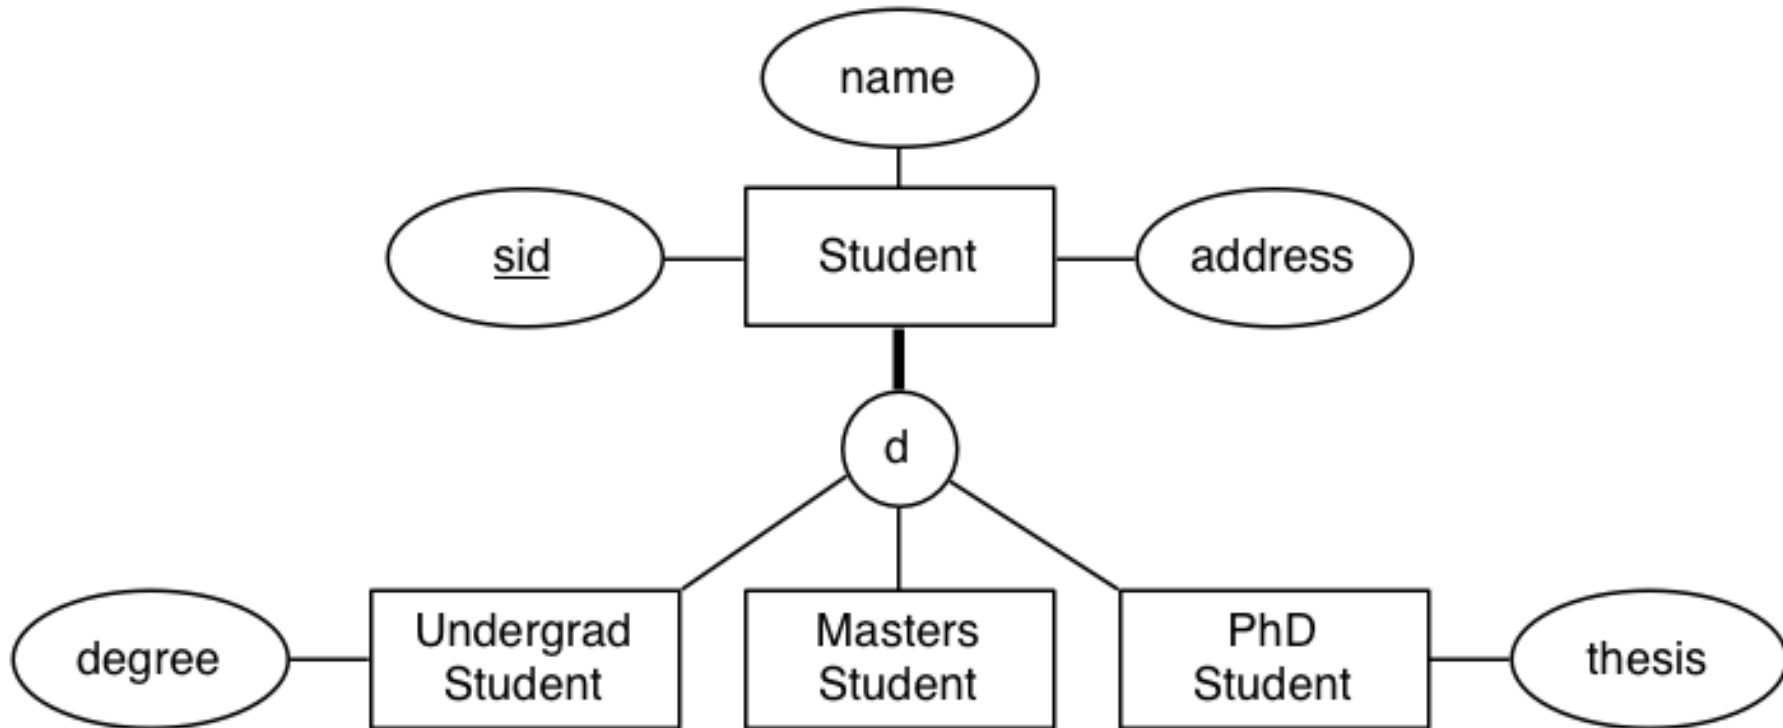
Person

<u>SSN</u>	name	address	salary	position	bonus
------------	------	---------	--------	----------	-------

NULL for Person who is not Employee

NULL for Employee who is not Manager

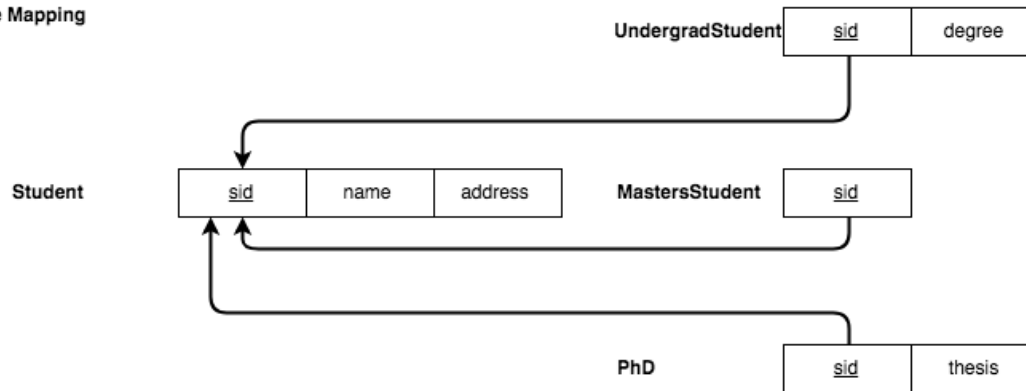
(5) Exercise: Disjoint subclasses



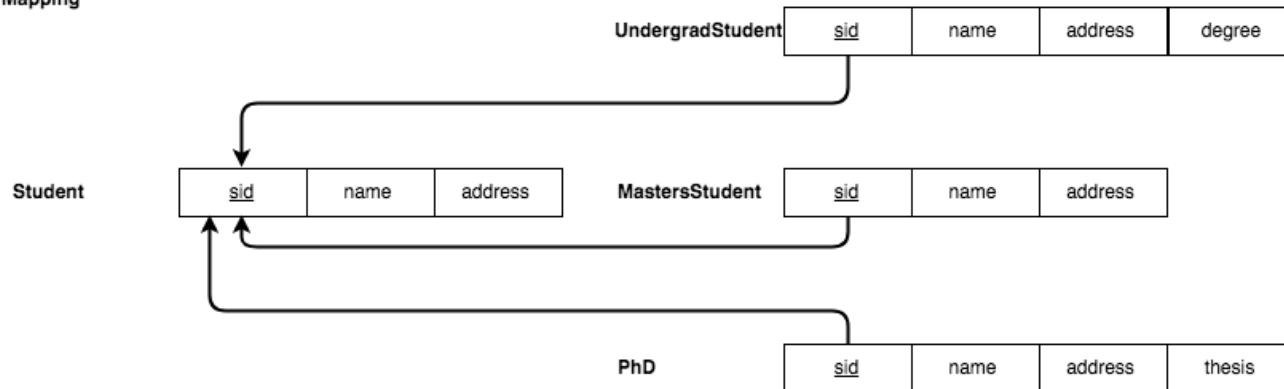
Use (a) ER-mapping, (b) OO-mapping, (c) 1-table-mapping

(5) Solution: Disjoint subclasses


ER Style Mapping



OO Style Mapping

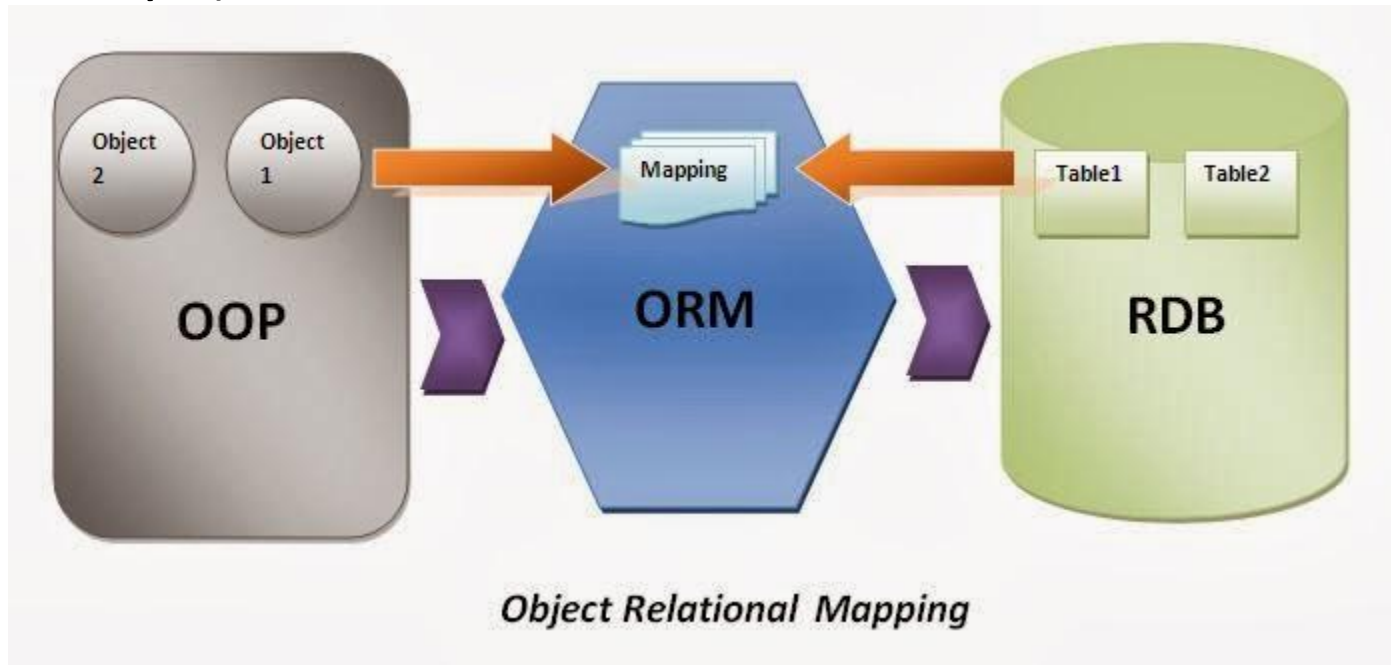


Single Table Style Mapping



Introduction to Object Relational Mapper (ORMs)

- A high-level abstraction framework that maps a relational database system to objects
- Automates all the CRUD (create/retrieve/update/delete) operations
- ORM is agnostic to which relational database is used (well at least theoretically...)



- Many different ORM frameworks around e.g., Hibernate, TopLink, SQLAlchemy

Why are ORMs useful?

- Shields the developer from having to write complex SQL statements and focus on the application logic using their choice of programming language
- Harmonisation of data types between the OO language and the SQL database
- Automates transfer of data stored in relational database tables into objects

Relational database (such as PostgreSQL or MySQL)

ID	FIRST_NAME	LAST_NAME	PHONE
1	John	Connor	+16105551234
2	Matt	Makai	+12025555689
3	Sarah	Smith	+19735554512
...

Python objects

```
class Person:  
    first_name = "John"  
    last_name = "Connor"  
    phone_number = "+16105551234"
```

```
class Person:  
    first_name = "Matt"  
    last_name = "Makai"  
    phone_number = "+12025555689"
```

```
class Person:  
    first_name = "Sarah"  
    last_name = "Smith"  
    phone_number = "+19735554512"
```

ORMs provide a bridge between
relational database tables, relationships
and fields and Python objects